

Applied Cryptography and Network Security

Adam J. Lee

adamlee@cs.pitt.edu

6111 Sennott Square

Lecture #6: Hash Functions

January 23, 2014




University of Pittsburgh



What is a hash function?

Definition: A **hash function** is a function that maps a **variable-length** input to a **fixed-length** code



Hash functions are sometimes called **message digest** functions

- SHA-1 stands for the **secure hash algorithm**
- MD5 stands for **message digest** algorithm (version 5)

In order to be useful cryptographically, a hash function needs to have a “randomized” output



For example:

- Given a large number of inputs, any given bit in the corresponding outputs should be set about half of the time
- Any given output should have half of its bits set on average
- Given two messages m and m' that are very closely related, $H(m)$ and $H(m')$ should appear completely uncorrelated

Informally: The output of an m -bit hash function should appear as if it was created by flipping m unbiased coins

Theoretical cryptographers sometimes use a more formalized notion of **random oracles** to model hash functions when analyzing security protocols



More formally, cryptographic hash functions should have the following three properties

Assume that we have a hash function $H : \{0,1\}^* \rightarrow \{0,1\}^m$

What does infeasible mean?!?

1. **Preimage resistance:** Given a hash output value z , it should be **infeasible** to calculate a message x such that $H(x) = z$
 - I.e., H is a **one way** function
 - Ideally, computing x from z should take $O(2^m)$ time
2. **Second preimage resistance:** Given a message x , it is infeasible to calculate a second message y such that $H(x) = H(y)$
 - Note that this attack is **always possible** given infinite time (**Why?**)
 - Ideally, this attack should take $O(2^m)$ time
3. **Collision resistance:** It is infeasible to find two messages x and y such that $H(x) = H(y)$
 - Ideally, this attack should take $O(2^{m/2})$ time

Why only $O(2^{m/2})$?



The Birthday Paradox!

The gist: If there are more than 23 people in a room, there is a better than 50% chance that two people have the same birthday

Wait, what?

- 366 possible birthdays
- **To solve:** Find probability p_n that n people all have *different* birthdays, then compute $1 - p_n$

$$p_n = \frac{365}{366} \frac{364}{366} \frac{363}{366} \cdots \frac{367 - n}{366}$$

- If $n = 22$, $1 - p_n \approx 0.475$
- If $n = 23$, $1 - p_n \approx 0.506$

Note: The value of n can be approximated as $1.1774 \times \sqrt{n} = 1.1774 \times \sqrt{366} \approx 22.525$

What the heck does this have to do with hash functions?!



Note that “birthday” is just a function $b : \text{person} \rightarrow \text{date}$

Goal: How many inputs x to the function b do we need to consider to find x_i, x_j such that $b(x_i) = b(x_j)$?


We're looking for collisions in the birthday function!

Now, a hash is a function $H : \{0, 1\}^* \rightarrow \{0, 1\}^m$

- **Note:** H has 2^m possible outputs

So, using our approximation from the last slide, we'd need to examine about $1.1774 \times \sqrt{2^m} = 1.1774 \times 2^{m/2} = O(2^{m/2})$ inputs to find a collision!

What are some things that we can do with a hash function?



Mutual Authentication



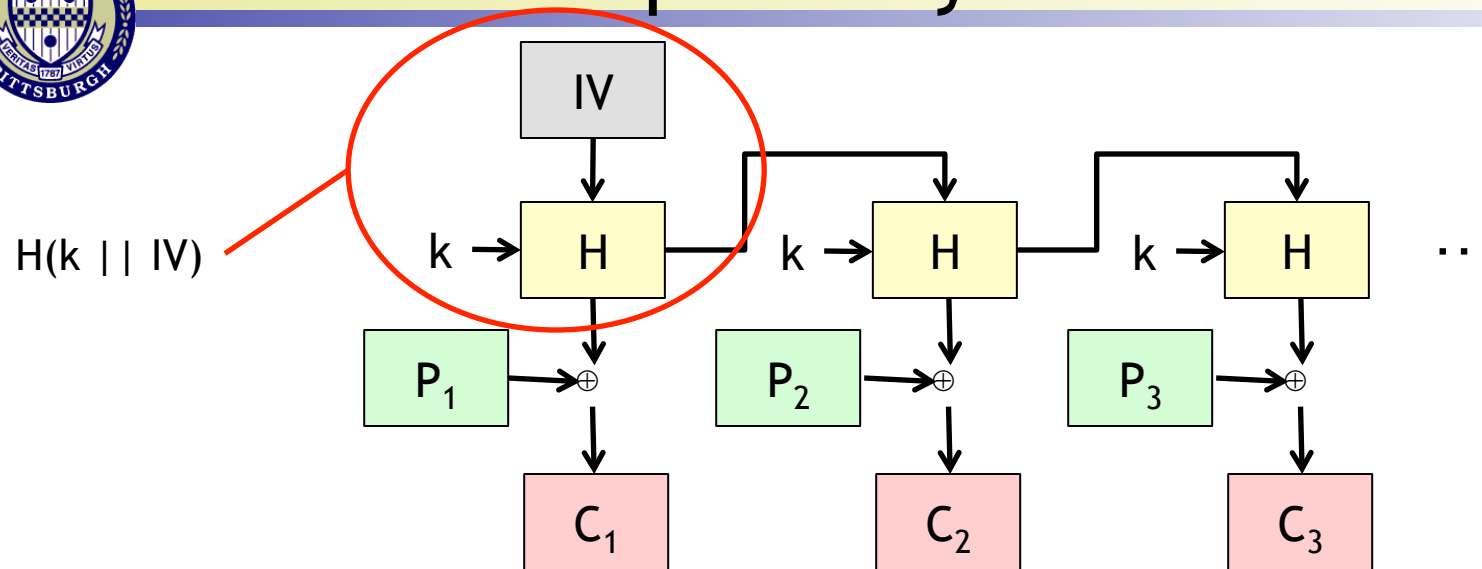
Document Fingerprinting

- Use $H(D)$ to see if D has been modified
- **Example:** Tripwire

MAC Functions

- Assume a shared key K
- Sender:
 - Compute $c = E_K(H(m))$
 - Transmit m and c
- Receiver:
 - Compute $c = E_K(H(m))$
 - Transmit m and c

Hash functions can even be used to generate cipher keystreams



Question: What block cipher mode does this remind you of?

- Output feedback mode (OFB), of course!

Why is this safe to do?

- Remember that hash functions need to have behave “randomly” in order to be used in cryptographic applications
- Even if the adversary knows the IV, he cannot figure out the keystream without also knowing the key, k



Hash functions also provide a means of safely storing user passwords

Consider the problem of safely logging into a computer system

Option 1: Store $\langle \text{username}, \text{password} \rangle$ pairs on disk

- **Correctness:** This approach will certainly work
- **Safety:** What if an adversary compromises the machine?
 - All passwords are leaked!
 - This probably means the adversary can log into your email, bank, etc...

Option 2: Store $\langle \text{username}, H(\text{password}) \rangle$ pairs on disk

- **Correctness:**
 - Host computes $H(\text{password})$
 - Checks to see if it is a match for the copy stored on disk
- **Safety:** Stealing the password file is less* of an issue

The previous applications provide us with an intuitive way to understand the importance of a hash function's cryptographic properties



1. **Preimage resistance:** Given a hash output value z , it should be infeasible to calculate a message x such that $H(x) = z$

Without this, we could recover hashed passwords!

2. **Second preimage resistance:** Given a message x , it is infeasible to calculate a second message y such that $H(x) = H(y)$

- **Example:** File integrity checking

- Say the `ls` program has a fingerprint f
- We could create a malicious version of `ls` that actually executes `rm -rf *`, but has the same document fingerprint

3. **Collision resistance:** It is infeasible to find two messages x and y such that $H(x) = H(y)$

In lecture 16, we'll see that this can lead to attacks that let us inject arbitrary content into protected documents!

Ok, enough high-level talk. How do these things actually work?



It is perhaps unsurprising that hash functions are effectively compression functions that are iterated many times

- **Compression:** Implied by the ability to map a large input to a small output
- **Iteration:** Helps “spread around” input perturbations

The book spends a lot of time talking about the “MD” family of message digest functions developed by Professor Ron Rivest (MIT)

Bad news: the most recent MD function, MD5, was broken in 2008

- Specifically, it has been shown possible to generate MD5 collisions in $O(2^{32})$ time, which is **much** faster than the theoretical “best case” of $O(2^{64})$
- We’ll talk more about this in Lecture 16

Rather than discuss MD5, we’ll focus on SHA-1

SHA-1 is built using the Merkle-Damgård construction

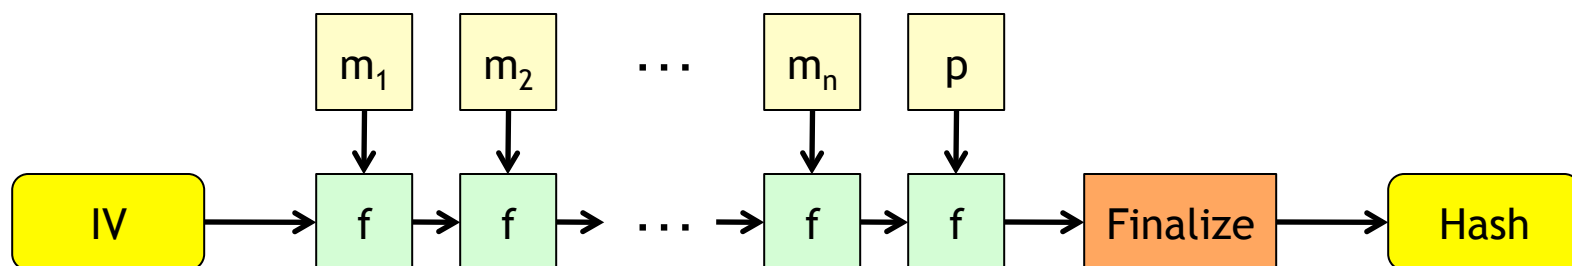


The **Merkle-Damgård construction** is a “template” for constructing cryptographic hash functions


- Proposed in the late '70s
- Named after Ralph Merkle and Ivan Damgård

Essentially, a Merkle-Damgård hash function does the following:

1. Pad the input message if necessary
2. Initialize the function with a (static) IV *← why is a static IV ok?*
3. Iterate over the message blocks, applying a **compression function** f
4. Finalize the hash block and output



Merkle and Damgård independently showed that the resulting hash function is **secure** if the compression function is **collision resistant**



Output: A 160-bit digest

- Pad message to a multiple of 512 bits
- Process one 512 bit chunk at a time
- Expand the sixteen 32-bit words into eighty 32-bit words
- Initialize five 32-bit words of state
- For each block of five 32-bit words
 - Apply function at right
 - Add result to output
- Concatenate five 32-bit words of output state





Initialization and Padding

Initialize variables:

$h0 = 0x67452301$

$h1 = 0xEFCDAB89$

$h2 = 0x98BADCFE$

$h3 = 0x10325476$

$h4 = 0xC3D2E1F0$


Note: These variables comprise the internal state of SHA-1. They are continuously updated by the compression function, and are used to construct the final 160-bit hash value.

Pre-processing:

append the bit '1' to the message

append $0 \leq k < 512$ bits '0', so that the resulting message length (in bits)
is congruent to $448 \equiv -64 \pmod{512}$

append length of message (before pre-processing), in bits, as 64-bit big-endian integer



Example:

$0xDEADBEEF \rightarrow 0xDEADBEEF8000 \dots 0020$

32 bits

$32_{10} = 0x20$



Initializing the compression function

Process the message in successive 512-bit chunks:

break message into 512-bit chunks

for each chunk

break chunk into sixteen 32-bit big-endian words $w[i]$, $0 \leq i \leq 15$

Extend the sixteen 32-bit words into eighty 32-bit words:

for i from 16 to 79

$w[i] = (w[i-3] \text{ xor } w[i-8] \text{ xor } w[i-14] \text{ xor } w[i-16]) \lll 1$

Initialize hash value for this chunk:

$a = h0$

$b = h1$

$c = h2$

$d = h3$

$e = h4$

Note: \lll denotes a left rotate.

Example: $00011000 \lll 4$


10000001



Main body of the compression function

Main loop:

for i from 0 to 79

if $0 \leq i \leq 19$ then

$f = (b \text{ and } c) \text{ or } ((\text{not } b) \text{ and } d); k = 0x5A827999$

else if $20 \leq i \leq 39$

$f = b \text{ xor } c \text{ xor } d; k = 0x6ED9EBA1$

else if $40 \leq i \leq 59$

$f = (b \text{ and } c) \text{ or } (b \text{ and } d) \text{ or } (c \text{ and } d); k = 0x8F1BBCDC$

else if $60 \leq i \leq 79$

$f = b \text{ xor } c \text{ xor } d; k = 0xCA62C1D6$

$\text{temp} = (a \lll 5) + f + e + k + w[i]$

$e = d; d = c; c = b \lll 30; b = a; a = \text{temp}$

Add this chunk's hash to result so far:

$h0 = h0 + a; h1 = h1 + b; h2 = h2 + c; h3 = h3 + d; h4 = h4 + e$

Note: Sometimes, we treat state as a bit vector...

... but other times, it is treated as an unsigned integer



Finalizing the result

Produce the final hash value (big-endian):

output = h0 || h1 || h2 || h3 || h4

"||" denotes concatenation

Interesting note:

- $k_1 = 0x5A827999 = 2^{30} \times \sqrt{2}$
- $k_2 = 0x6ED9EBA1 = 2^{30} \times \sqrt{3}$
- $k_3 = 0x8F1BBCDC = 2^{30} \times \sqrt{5}$
- $k_4 = 0xCA62C1D6 = 2^{30} \times \sqrt{10}$

Question: Why might it make sense to choose the k values for SHA-1 in this manner?



SHA-1 in Practice

SHA-1 has fairly good randomness properties

- SHA1("The quick brown fox jumps over the lazy dog")
 - 2fd4e1c6 7a2d28fc ed849ee1 bb76e739 1b93eb12
- SHA1("The quick brown fox jumps over the lazy **c**og")
 - de9f2c7f d25e1b3a fad3e85a 0bd17d9b 100db4b3

In the above example, changing 1 character of input alters 81 of the 160 bits in the output!

To date, the best attack on SHA-1 can find a collision with about $O(2^{61})$ steps; in theory, this attack *should* take $O(2^{80})$ steps.

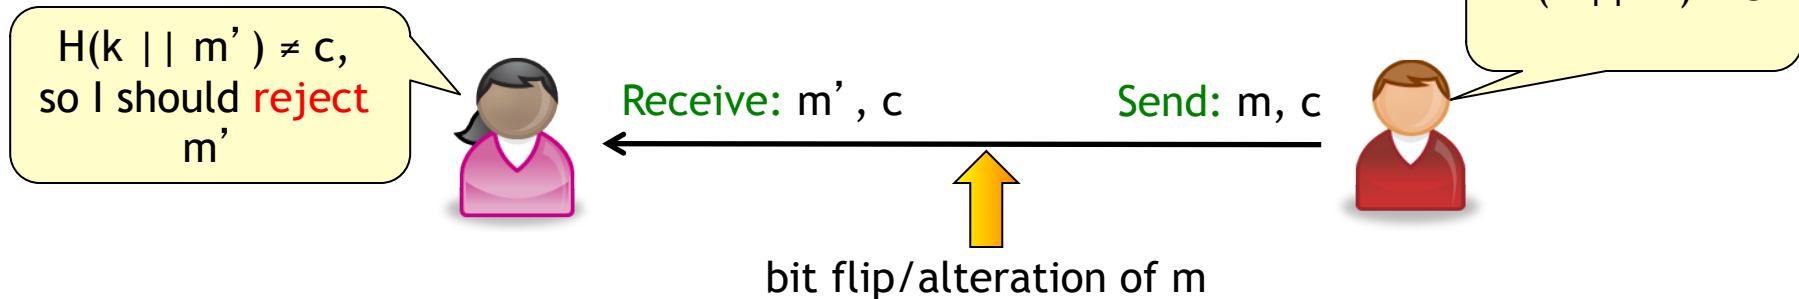
As a result, NIST ran a hash function competition to design a replacement for SHA-1 (Keccak chosen in Oct 2012)

Like the AES competition

Although hashes are unkeyed functions, they can be used to generate MACs




A keyed hash can be used to detect errors in a message



Unfortunately, this isn't *totally* secure... (Why?)

- It's usually easy to add more data while still generating a correct MAC!



There are also attacks against $H(m || k)$ and $H(k || m || k)$!

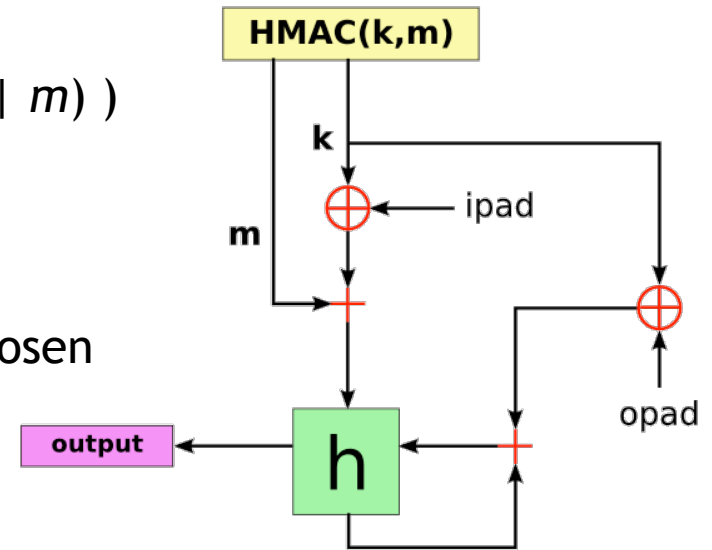


HMAC is a construction that uses a hash function to generate a **cryptographically strong** MAC

$$\text{HMAC}(k, m) = H((k \oplus \text{opad}) || H((k \oplus \text{ipad}) || m))$$

- opad = 01011010
- ipad = 00110110

The opad and ipad constants were carefully chosen to ensure that the internal keys have a large **Hamming distance** between them



Note that H can be **any** hash function. For example, HMAC-SHA-1 is the name of the HMAC function built using the SHA-1 hash function.


Benefits of HMAC:

- Hash functions are faster than block ciphers
- Good security properties (**Why?**)
- Since HMAC is based on an **unkeyed** primitive, it is not controlled by export restrictions!



Hash functions can also help us check the integrity of large files efficiently

Many peer-to-peer file sharing systems use **Merkle trees** for this purpose



Why is this good?

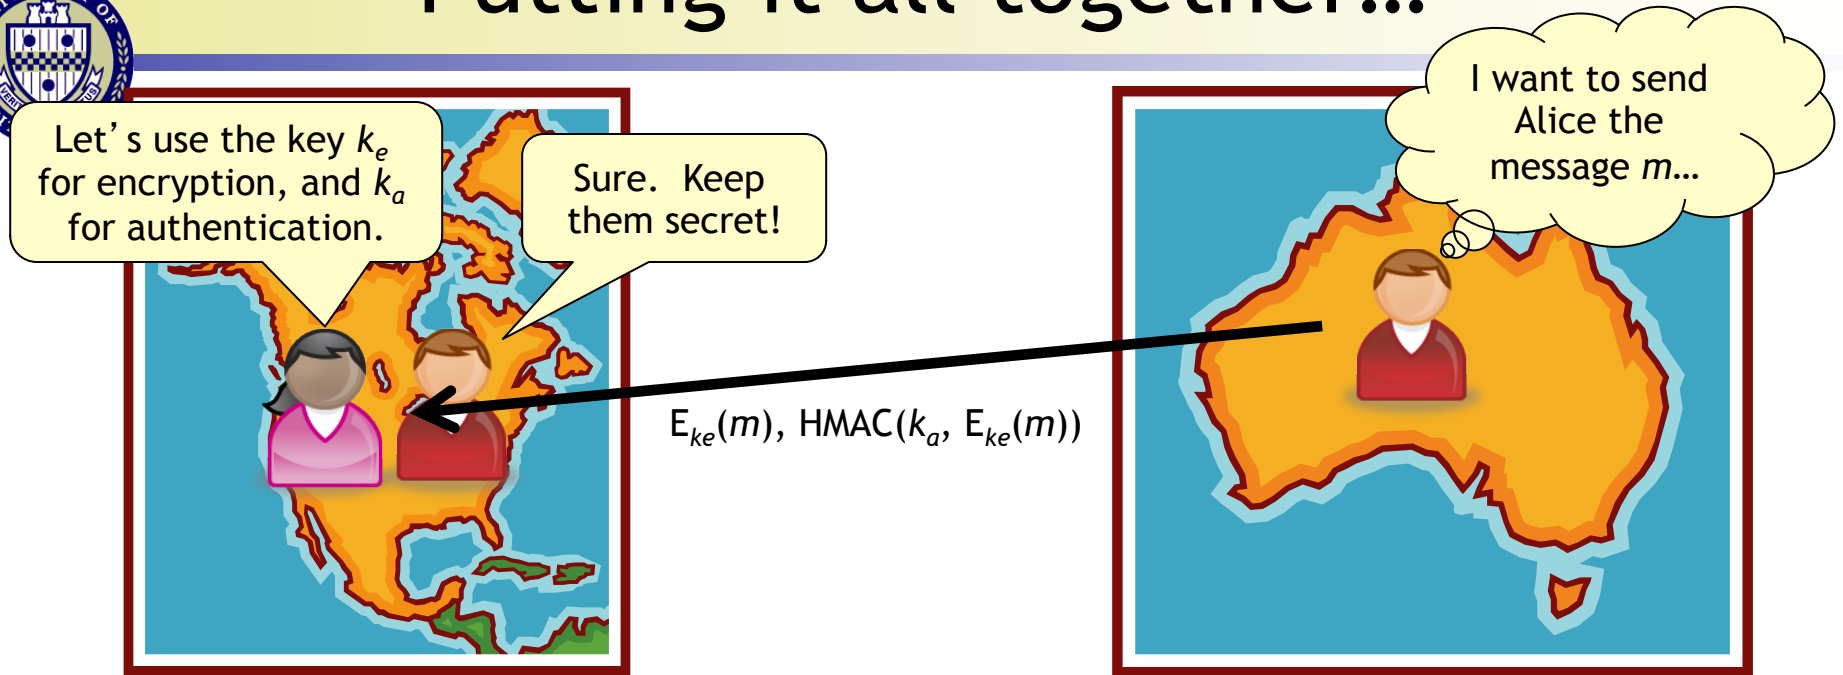
- One branch of the hash tree can be downloaded and verified at a time
- Interleave acquisition of integrity check with file data
- Errors can be corrected on the fly

BitTorrent uses **hash lists** for file integrity verification

- Must download full hash list prior to verification



Putting it all together...



Why compute the HMAC over $E_{k_e}(m)$?

- Alice doesn't need to waste time decrypting m if it was mangled in transit, since its authenticity can be checked first!

Why use two separate keys?

- In general, it's a bad idea to use cryptographic material for multiple purposes



Project #2