

Applied Cryptography and Network Security

Adam J. Lee

adamlee@cs.pitt.edu

6111 Sennott Square

Lecture #11: Handshake Protocols

February 11, 2014



University of Pittsburgh




Announcements

Project Phase 3 goes out tonight

HW 1 is due next week



Poll Results



Most common challenging topics:

- AES details and RSA math
- Acronyms



Today's Topics

Today we'll start looking at four types of handshake protocols:

- Login only protocols
- Mutual authentication protocols
- Integrity/encryption setup protocols
- Mediated authentication protocols

As we'll see, there is a lot of subtlety that goes into designing these types of protocols

In short, this is a detail-oriented lecture...



Notation Overview

Types of keys:

- K_{AB} : a secret key shared between A and B
- k_A : The public key belonging to A
- k_A^{-1} : The private key belonging to A

Types of cryptographic operations

- $\{ M \}_{K_{AB}}$: Message M encrypted with the secret key K_{AB}
- $\{ M \}_{k_A}$: Message M encrypted using A's public key
- $[M]_{k_A^{-1}}$: Message M signed using A's private key

Miscellany:


- R_A : A random number chosen by A
- $\{ M || R_A \}$: The concatenation of M and R_A



Login only protocols



Login-only protocols are designed to authenticate the user prior to permitting system access



This protocol operates under the assumptions that:

- Only the user knows his password
- The password database is private
- No one is listening on the communication channel


Question: Are these assumptions always valid? Always invalid?

- It depends!

Fortunately, this botched protocol is easy to (partially) fix




It is natural to turn this basic protocol into a cryptographic challenge/response using a shared secret



Interesting notes:

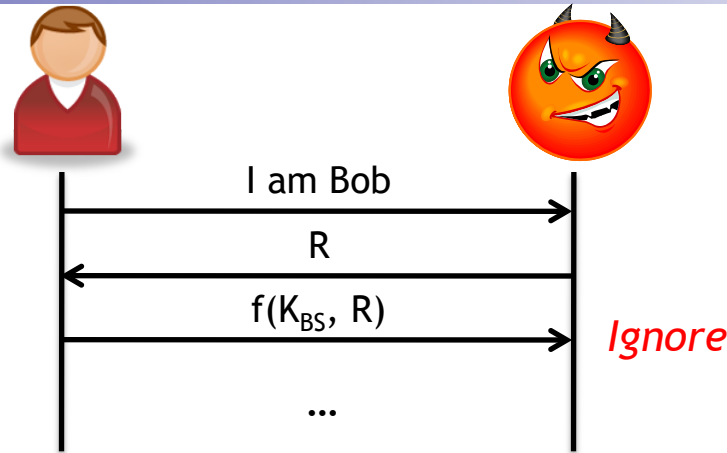
- The use of K_{BS} convinces the server that he is talking to Bob
- f can be either a two-way (encryption) or one-way (hash) function



Note: We **must** use a two-way (encryption) function (**Why?**)

Unfortunately, these protocols are subject to a variety of attacks...

Since Bob does not authenticate the server, anyone can pose as the server!



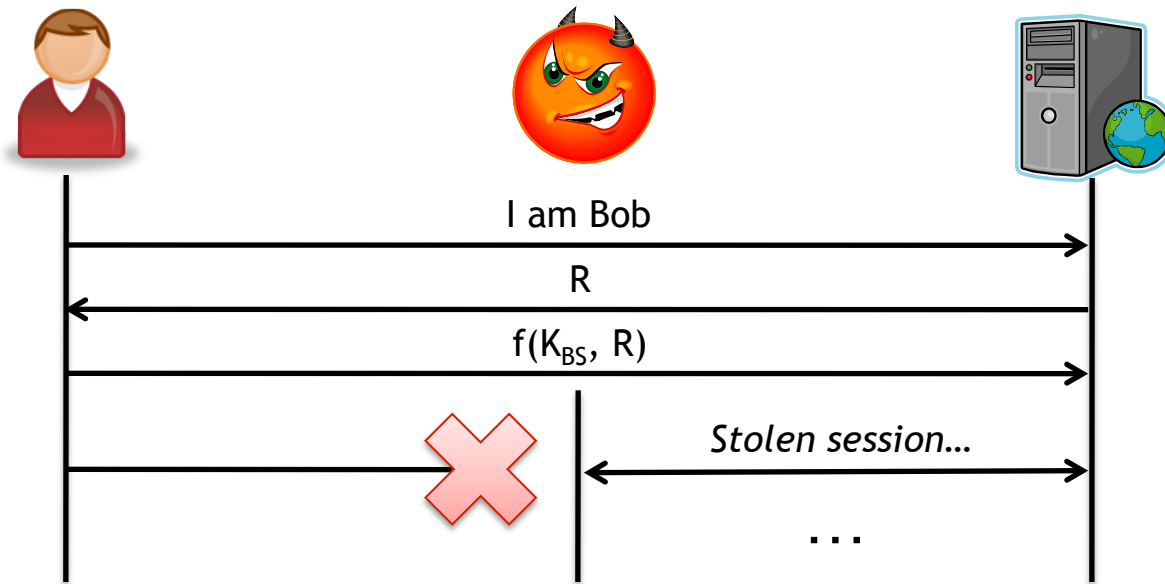
Question: Why might this be a problem?

If the attacker has any knowledge of the server, he can:

- Steal Bob's password(s) as he logs in to, e.g., his email
- Use $f(K_{BS}, R)$ to act as a **man in the middle**
- If K_{BS} is derived from a password, $f(K_{BS}, R)$ can be used to launch an **offline** password guessing attack



An attacker capable of blocking traffic from Bob can hijack Bob's session after login



Who could this attacker be?


- Malicious router
- Peer on same network
- ...

*The protocol ends
after the user is
authenticated!*

Question: Why is this attack possible?



What if we only have one message in the protocol to work with?




This protocol has several **strong** points:

- Easily replaces protocols that rely on simply sending a cleartext password
- Only requires one message, not three
- The server does not need to maintain volatile state (e.g., previously used Rs)

Sadly, it also has some **weaknesses**:

- Bob and the server need synchronized clocks!
- Attackers snooping on the wire can reuse Bob's encrypted timestamp to log into other servers within an acceptable window of time
- If an attacker can convince the server to roll back its clock, old encrypted timestamps can be reused!

These protocols can also be adapted to use public key cryptography



Why do these protocols work?

- Protocol 1: Only Bob can generate $[R]_{k_B^{-1}}$
- Protocol 2: Only Bob can open $\{R\}_{k_B}$

Interesting note: No more sensitive databases of user passwords or shared secrets!

What is the problem with these protocols?

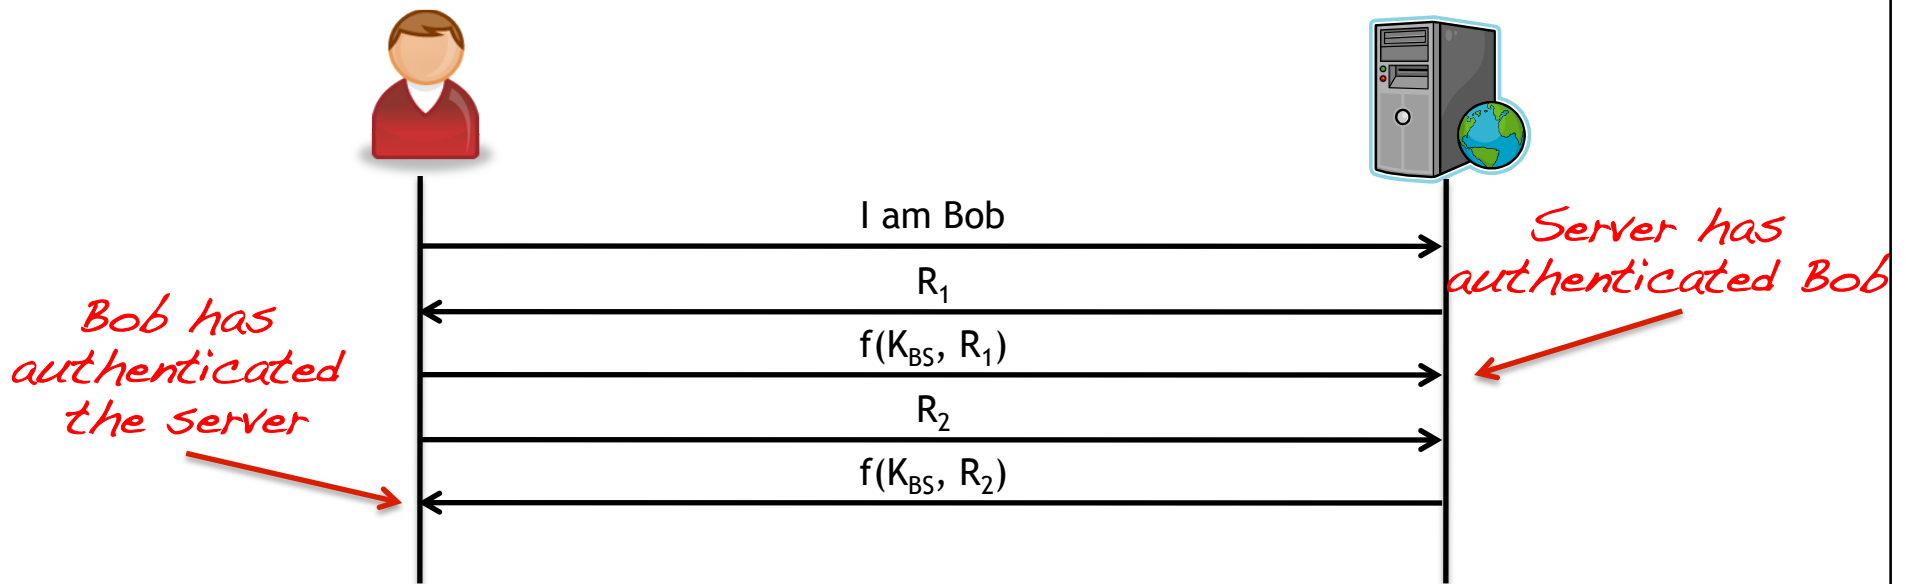


Mutual Authentication

Often times, both participants in a protocol want to authenticate one another




One way to do this is to (essentially) run two invocations of our earlier protocols:



This seems like a lot of messages, doesn't it? Can't we optimize this thing somehow?



A stab at optimization...



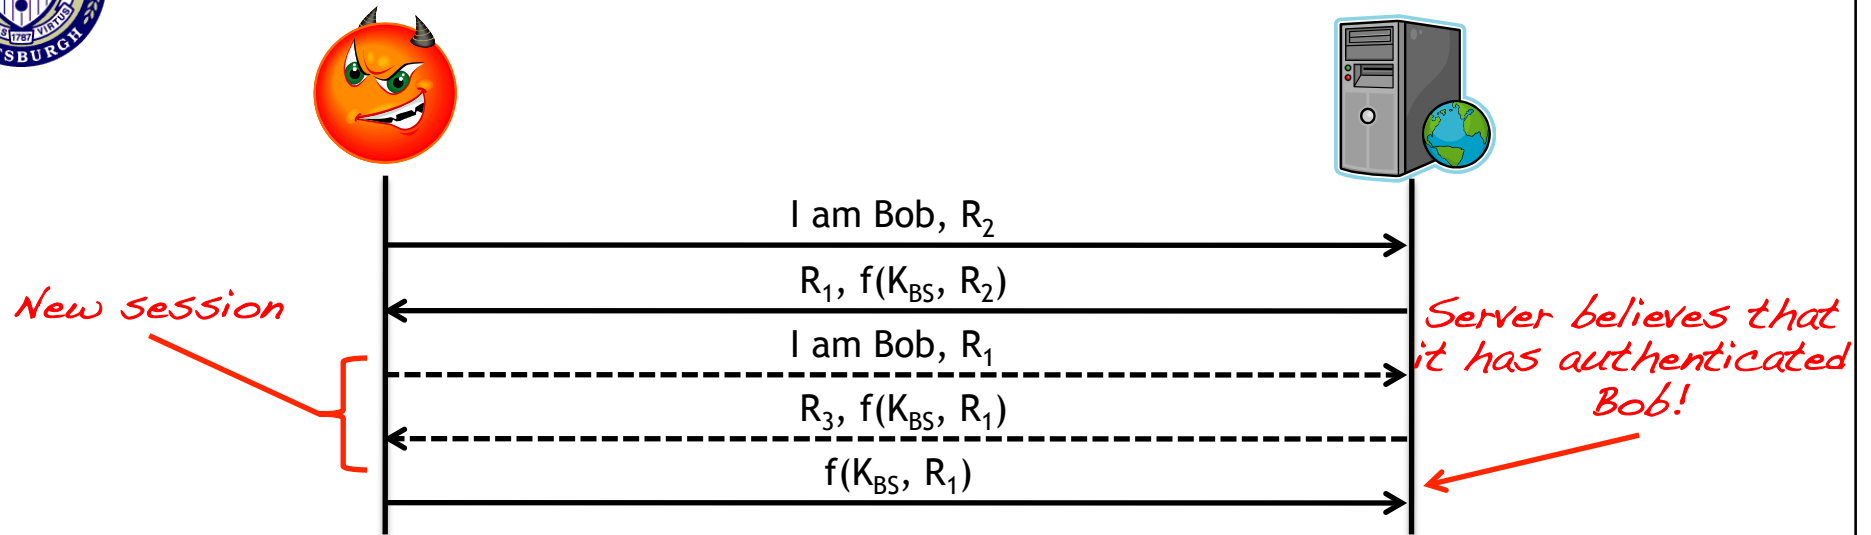
In theory:

- Bob has authenticated the server after message 2
- The server has authenticated Bob after message 3

In practice: This isn't actually the case!

This protocol is vulnerable to what is known as a **reflection attack**

This opens the door to what is known as a reflection attack



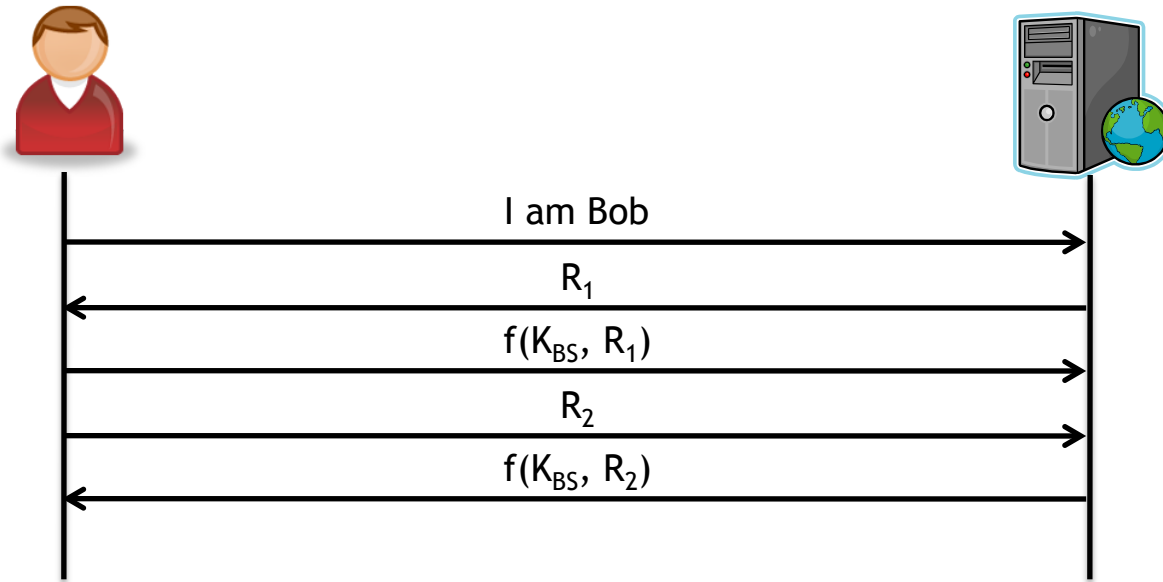
Why is this attack possible?

- Both parties do **exactly** the same thing
- No way to prevent “reflecting” challenges back

However, this weakness is easily avoided

- Use different keys in each direction
 - E.g., K_{BS} could be XORed with different constants in each direction
- Force challenges to encode direction of transmission
 - E.g., use ($\langle \text{sender name} \rangle || R$) instead of R alone

Why isn't the long protocol vulnerable to this attack?

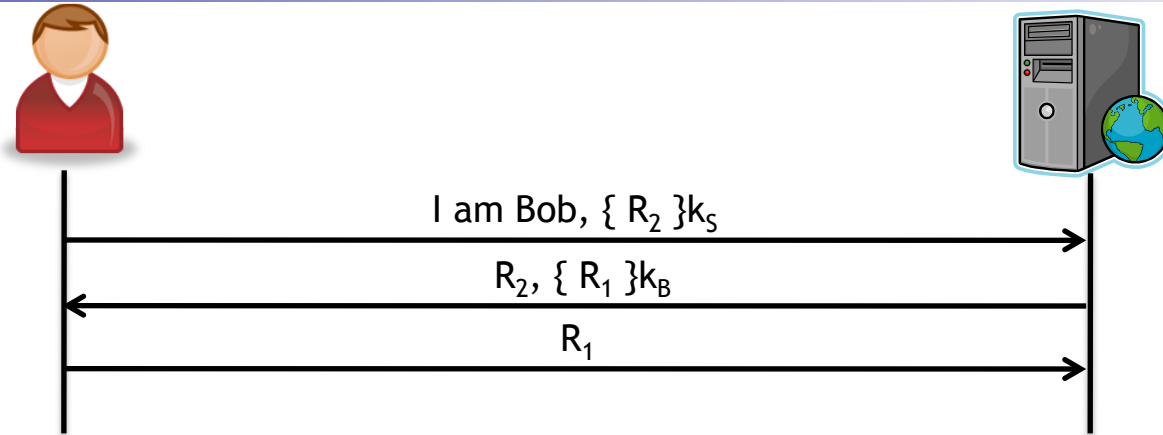


Answer: Bob needs to prove his identity to the server prior to the server proving its identity to Bob!

Lesson: Be careful when "optimizing" security protocols...



Mutual authentication protocols can also be constructed using public key cryptography



This looks a lot like our “optimized” secret key mutual authentication protocol, doesn’t it?

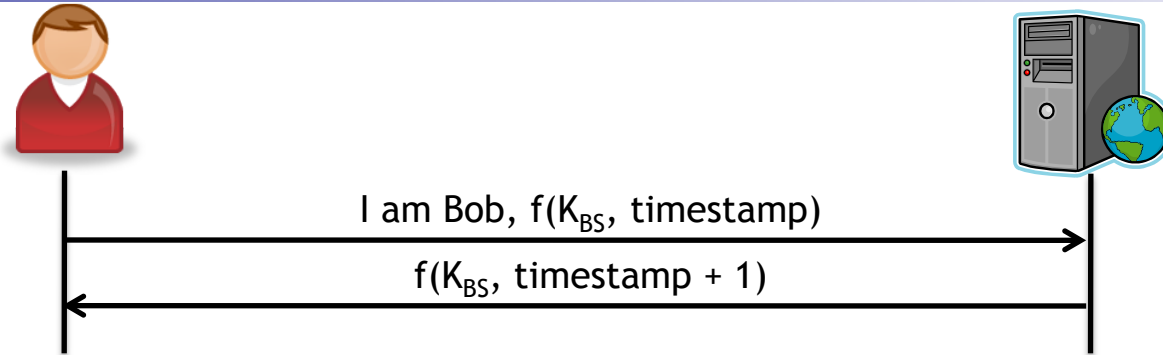
Question: Does this protocol suffer from the reflection attack, too?

Note: This protocol assumes that Bob and the server know each others’ public keys a priori. This could be done via:

- Acquisition from a trusted certificate authority
- Encrypt public keys with a symmetric key derived from a shared password
- Offline configuration (think SSH)



We can further reduce the number of messages by using timestamps instead of random challenges



This protocol is nice, as it fits within a two message exchange

- Request/reply exchanges
- RPC invocations
- Etc.

Question: Why do we need timestamps for this to work?

- Freshness! (Provided that clocks are **synchronized**)

Question: Why does the server return $f(K_{BS}, \text{timestamp} + 1)$?

- Protection against replay attacks!



Integrity/Encryption Setup Protocols



Often times, authentication is just not enough...

After authenticating, it is often necessary to protect the integrity and/or confidentiality of the rest of the conversation

Example: Telnet versus SSH

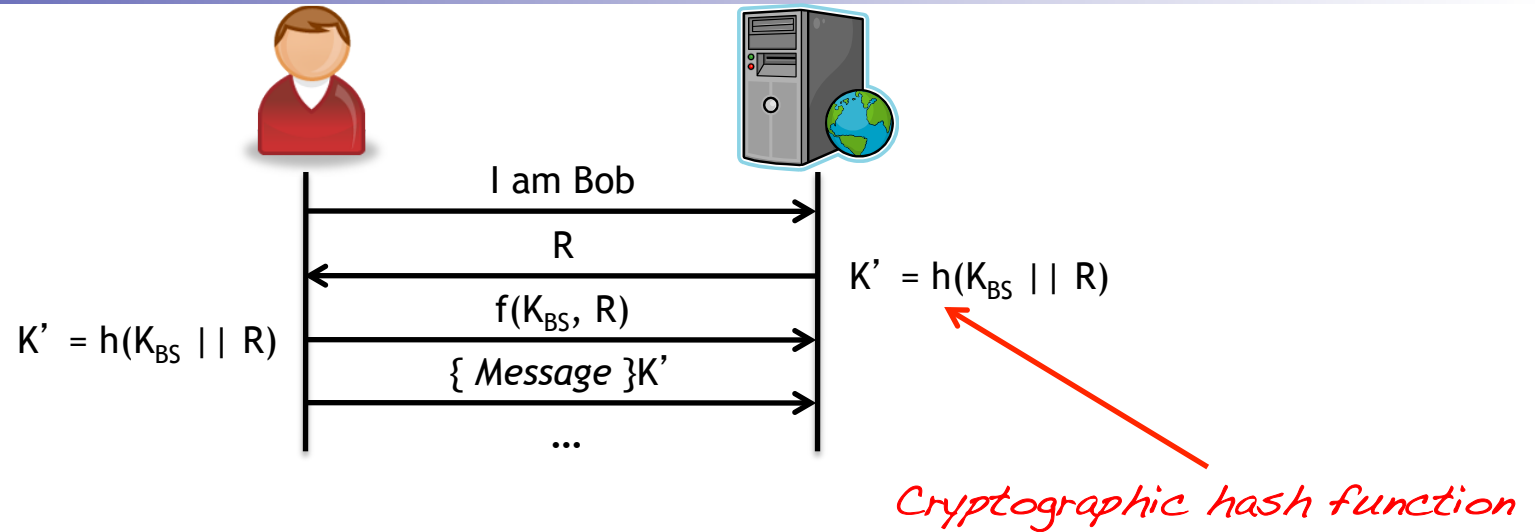
But wait, don't all of the authentication protocols that we've talked about require shared keys anyway?!?

It turns out that generating **fresh** keys regularly is important

- Overuse can make long term secret keys easier to break
- Per-session keys limit replay/injection attacks to a single session
- **Forward secrecy** of individual sessions
- ...



This turns out to be fairly simple to do!



Interesting notes:

- This protocol does not add **any** messages to our authentication exchange!
- Even though R is visible to the adversary, K' cannot be guessed (**Why?**)
- If the session key K' is leaked, the long term secret K_{BS} is still safe (**Why?**)

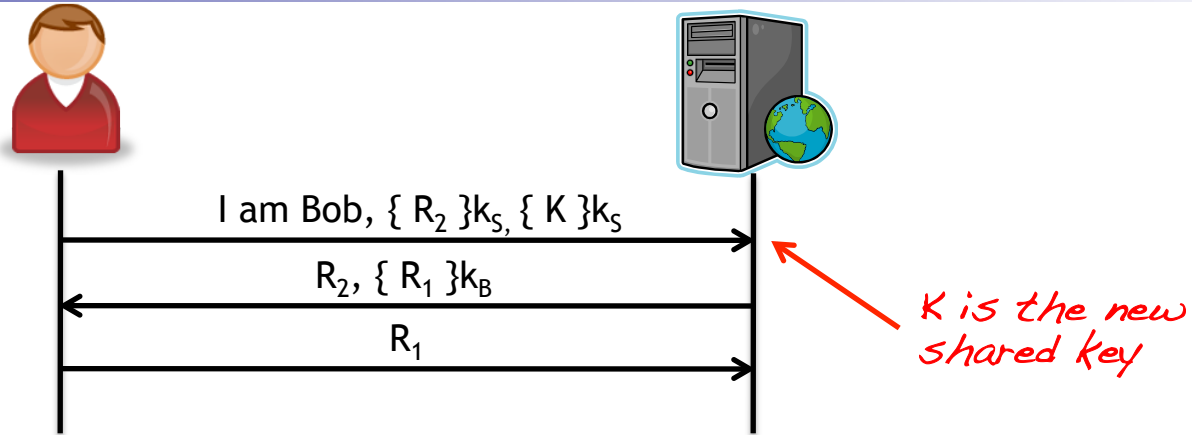
Mutual authentication protocols can be adapted in a similar manner

- E.g., $K' = h(K_{BS} || R_1 || R_2)$
- Note that order of R s is important!

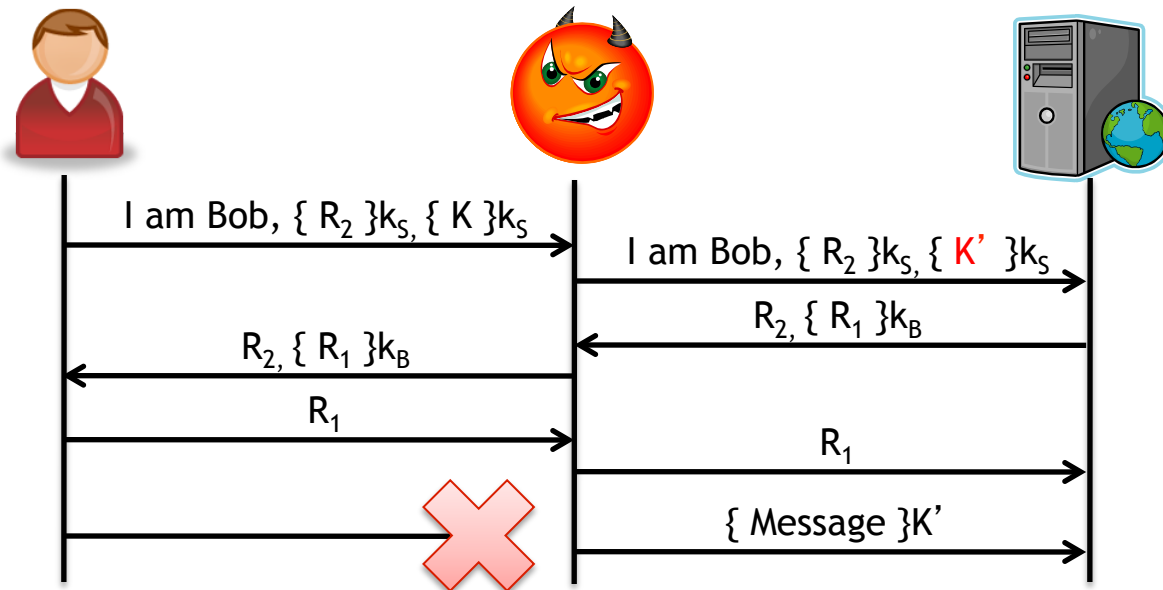
We can also derive session keys using public-key authentication protocols



Initial attempt:

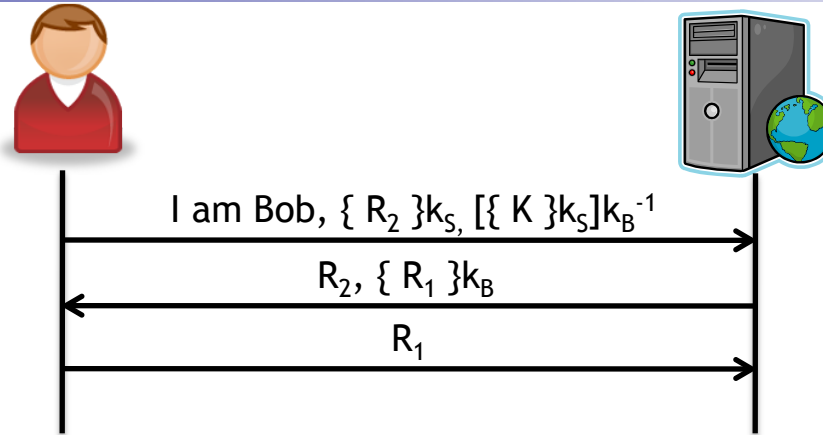


Unfortunately, this protocol can be **hijacked!**





Digital signatures can help us fix this problem!



Why does this work?

- The signature ensures that the key was actually generated by Bob
- This provides a **binding** between the authentication protocol and the key exchange protocol

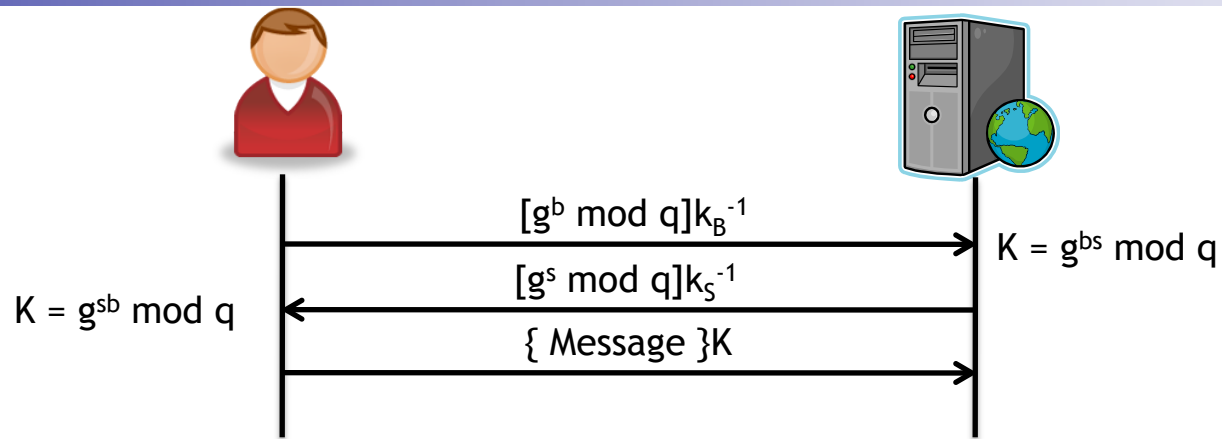
However, there are still issues with this protocol...

Assume that the server is eventually compromised

- This means that the adversary learns k_S^{-1}
- If the adversary recorded the above exchange, K can be recovered!



A signed Diffie-Hellman key exchange prevents this problem



As we learned earlier, the Diffie-Hellman exchange allows Bob and the server to agree on a shared secret key over a public channel

Even if both parties are later compromised (i.e., k_B^{-1} and k_S^{-1} are revealed) data encrypted with K is safe! (Why?)

Question: Why are the digital signatures needed?

What happens if we need to generate more than one shared key?



The previous protocols only provide us with a single shared key. How can we derive multiple keys?

One method is to run these protocols multiple times

- This is expensive for the participants
 - Fortunately, this is unnecessary!
-

Case study: Key derivation in SSH

- The SSH protocol uses a Diffie-Hellman exchange, and computes
 - A shared key K
 - An exchange hash value H
- Client to server encryption key: $h(K \parallel H \parallel \text{"C"} \parallel \text{session_id})$
- Client to server integrity key: $h(K \parallel H \parallel \text{"E"} \parallel \text{session_id})$

Question: Why is the above safe to do?



Mediated Authentication Protocols

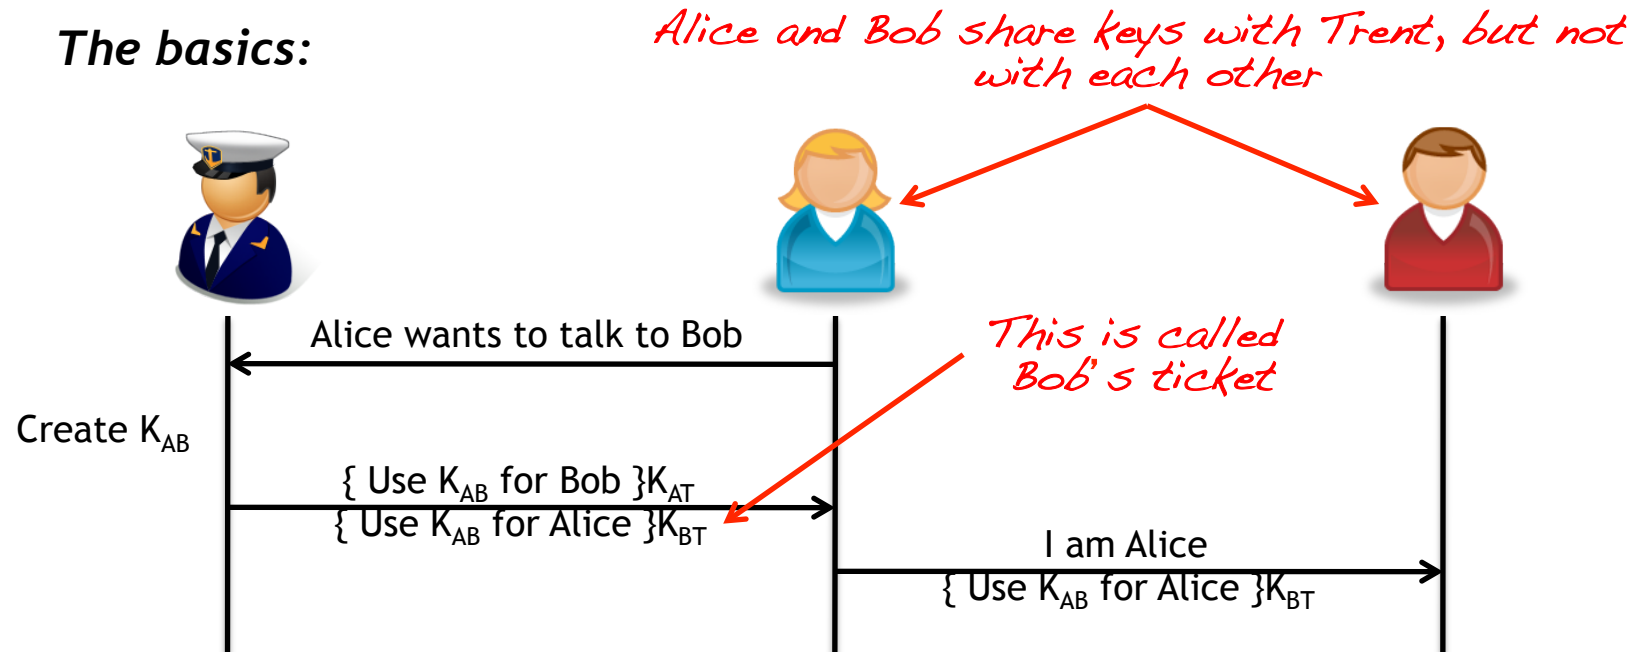
What if we like the speed of symmetric key cryptography, but not the key management headaches?



Specifically, how can we let two users securely establish a shared key?

Mediated authentication protocols make use of a trusted mediator, or key distribution center (KDC), to make this possible!

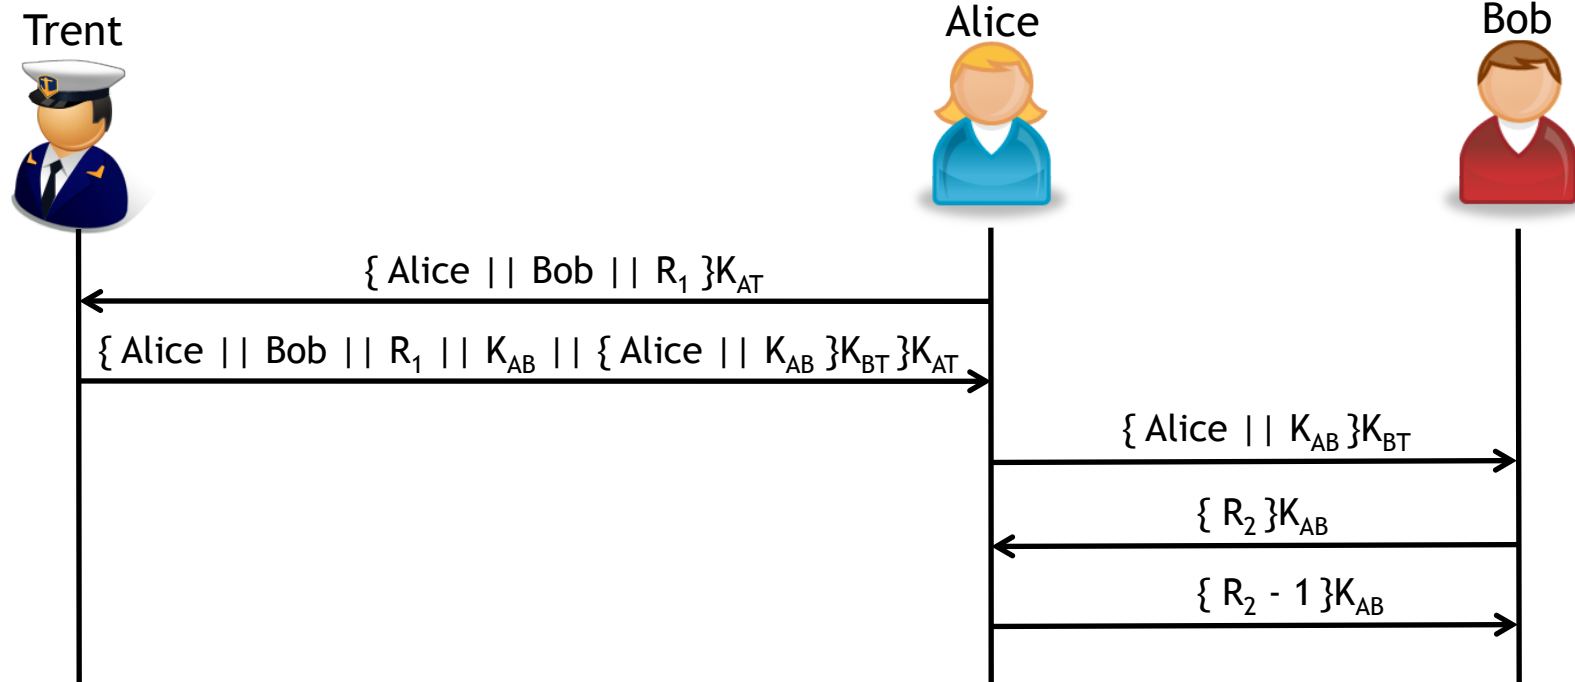
The basics:



Note: This protocol is incomplete, as it does not authenticate Alice and Bob to one another. However, it is a good place to start...



The Needham-Schroeder protocol is a well-known mediated authentication protocol

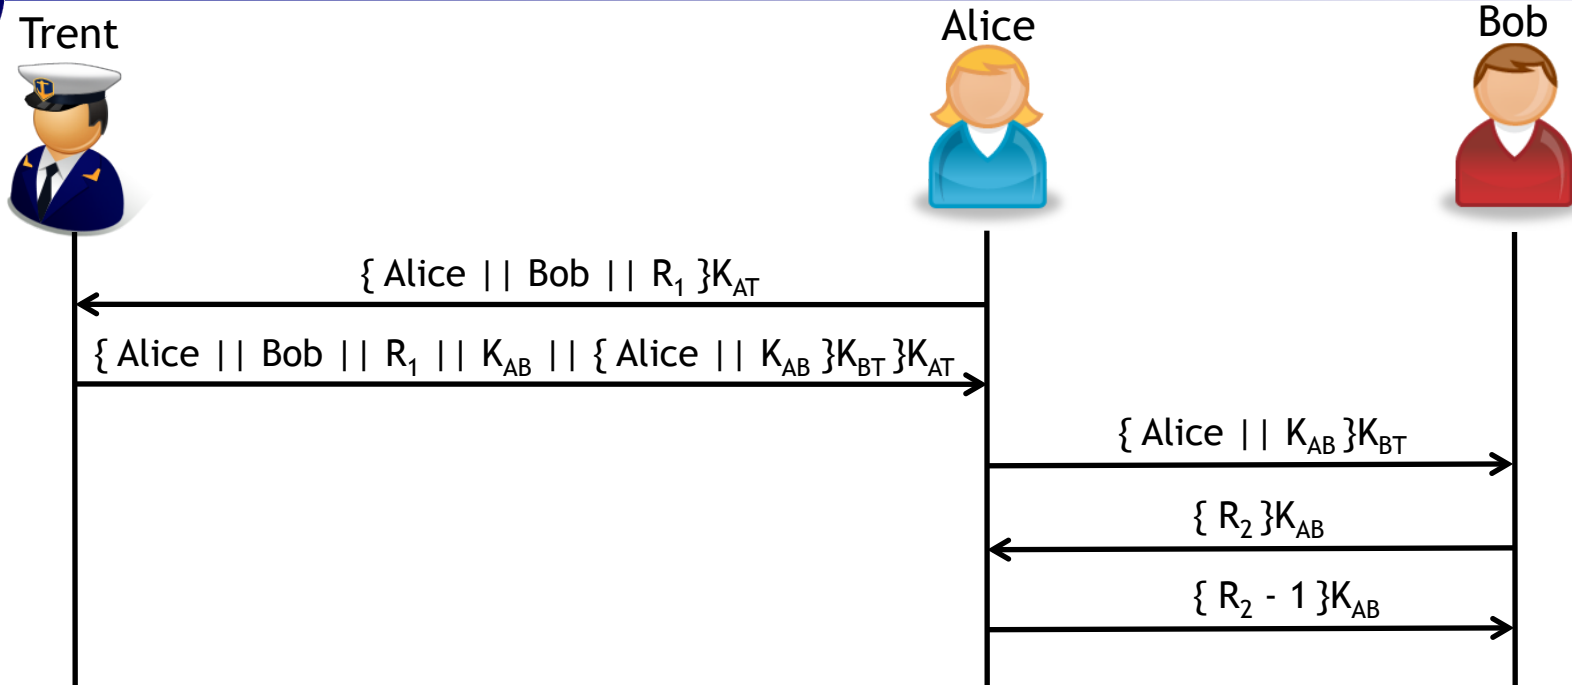


Note: R_1 and R_2 are called **nonces**

- Must be generated at random (unpredictable)
- Cannot be used in more than one protocol execution



Why does Needham-Schroeder work?



After message 2 Alice

- Knows that this message is fresh
- Knows that the session key is to be shared with Bob

After message 3, Bob knows that he has a shared key with Alice

After message 5, Bob knows that this key is fresh (Why?)

The Needham-Schroeder protocol assumes that all keys remain secret



Assume that Eve intercepts the message $\{ \text{Alice} \parallel K_{AB} \}_{K_{BT}}$ and later learns the **session key** K_{AB}

Question: *Why might a session key become compromised?*

Eve can now launch a replay attack!

- Eve can replay the message $\{ \text{Alice} \parallel K_{AB} \}_{K_{BT}}$
- She can intercept Bob's response $\{ R_3 \}_{K_{AB}}$ to Alice
- Since Eve knows K_{AB} , she can decrypt this message and reply $\{ R_3 - 1 \}_{K_{AB}}$

How can we defend against this type of attack?



The Ottway-Rees protocol prevents this attack

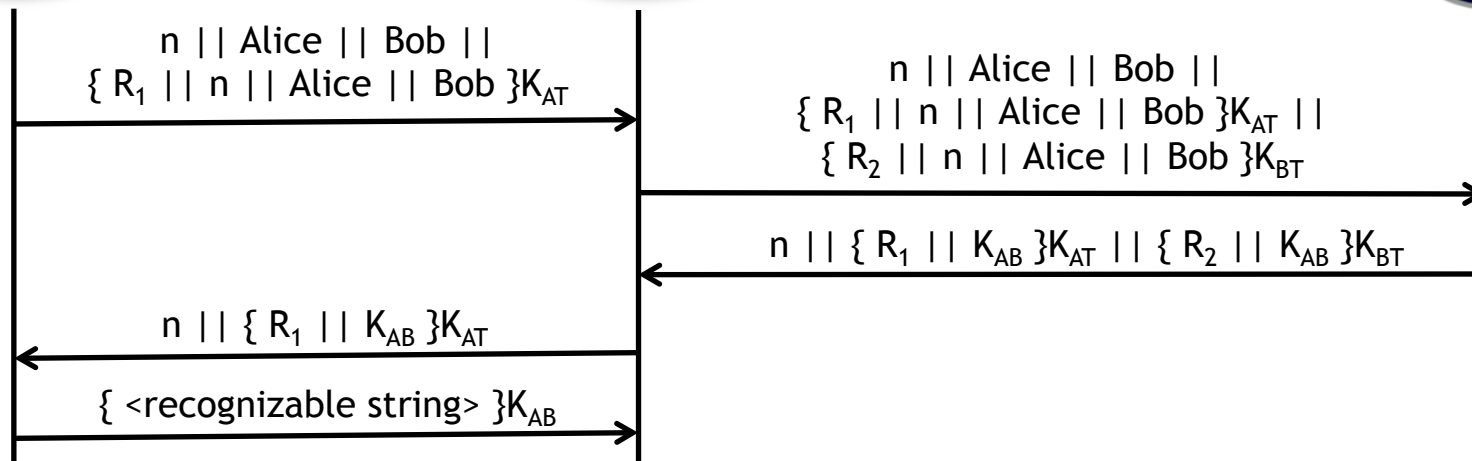
Alice



Bob



Trent



Properties of this protocol:

- After message 3, Bob knows that he has a **fresh** session key to share with Alice that was generated by Trent
- After message 4, Alice knows that she has a **fresh** session key to share with Bob that was generated by Trent
- After message 5, Bob knows that Alice received the shared key

Why doesn't a compromised session key subvert this protocol?



Assume that Eve:

- Records the message $n || \{ R_1 || K_{AB} \}_{K_{AT}} || \{ R_2 || K_{AB} \}_{K_{BT}}$
- Breaks the key K_{AB}

Now, say that Eve tries to forge a version of message 4 to Alice

- $n || \{ R_1 || K_{AB} \}_{K_{AT}}$

If Alice **does not** have an ongoing exchange with Bob, this forgery will fail

- Why? No initial state saved

If Alice **does** have an ongoing exchange with Bob

- The forgery will fail if the number n does not match
- If n does match, the forgery will fail because Alice will be using a different nonce R_1



Conclusions

So far, we've learned about four types of handshake protocols:

- Login only protocols
- Mutual authentication protocols
- Integrity/encryption setup protocols
- Mediated authentication protocols

These protocols are very simple, but very sensitive to change

Understanding the types of attacks that these protocols can be subjected to is a very important facet of designing secure networked systems

Next: Strong password protocols