

Applied Cryptography and Network Security

Adam J. Lee

adamlee@cs.pitt.edu

6111 Sennott Square

Lecture #23: OS and Application Security

April 3, 2014



University of Pittsburgh



Overview

OS protection model overview

- Memory protection
- Transitions between user mode and kernel mode

Integrity of mechanism

Compromising integrity of mechanism by violating assumptions at

- The OS level: Buffer overflow case study
- The application level: SQL injection case study



Protection in x86 Processors

x86 processors operate in two distinct modes

- Kernel mode (privileged)
- User mode (unprivileged)

The kernel is typically “small” when compared to the entire OS

- Controls base system functionality
- Provides interface to hardware
- Manages system data structures

Why use a small **trusted** kernel?

- In theory, a small kernel can be rigorously tested and verified
- In practice, it's always good to minimize your trusted computing base
 - Recall Saltzer and Schroeder's **principle of least privilege**

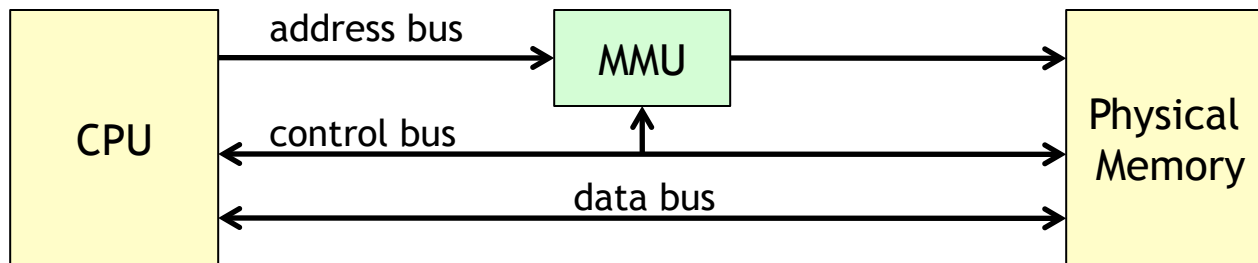


One critical job of the OS is memory protection

A user should be **unable** to access memory used by the kernel or other users

- This provides isolation, which is necessary to ensure integrity

How is this enforced?



The memory management unit (MMU) plays the role of the enforcer

- Translates logical memory addresses into physical addresses
- Allows regions of memory to be marked as read only
- Ensures that requests are a valid
 - Point to valid physical addresses
 - Are permissible (e.g., not writing to read only memory)

The OS controls the MMU and manages its lookup tables



It is unreasonable to believe that we can completely separate users from the kernel

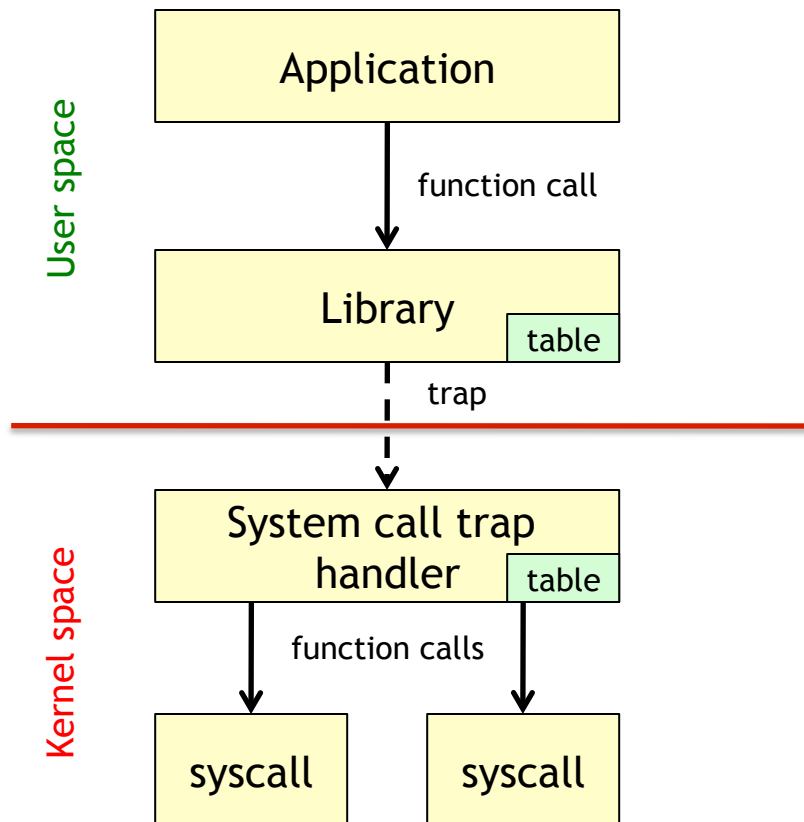
System calls allow user space code to access OS services

System calls **transfer control** of the CPU from the user to the kernel

How does this process work?

- User code identifies which kernel service it wants to invoke (e.g., by leaving a code in a specific register)
- Arguments for this service are left on the program stack
- Trap instruction issued
- Trap handler looks up code and gathers any arguments from the stack
- A function call is made

Recall from your OS class that the above process is called a **context switch**





OS security depends on integrity of mechanism

The protection mechanisms used by the OS are designed to narrow the interface through which users can access protected data

In order to ensure that the system remains secure, we need to ensure that

- The principle of **complete mediation** is followed
- “Trusted” code is actually trusted

There is also an **implicit assumption** that user processes behave as expected by the users who invoke them

As a result, compromising the security of an OS typically involves **violating** the **integrity** of its enforcement mechanisms





How can we violate the OS's assumption of integrity of mechanism?

Trivial attack: Extend the TCB by installing a malicious device driver

- E.g., Why not write a keyboard driver that's also keystroke logger?
- Drivers are trusted, and have complete access to kernel data structures
- This is a viable attack method, but its uses are limited


A better idea is to make an **existing** trusted process behave badly

Approach 1: Kernel Level



Use a buffer overflow to subvert the kernel's context switch process

Approach 2: Application Level



Violate application-level constraints via SQL injection



Buffer overflow is one of the oldest tricks in the book

In fact, buffer overflow was one of the attack vectors exploited by the Morris worm back in 1988!

We'll learn more about the Morris worm in later lectures

Informally, a buffer overflow works by providing an extremely long input to a user program that causes system data structures to be overwritten

The result of this is that control is transferred to code injected by the attacker, instead of being transferred back to the parent process

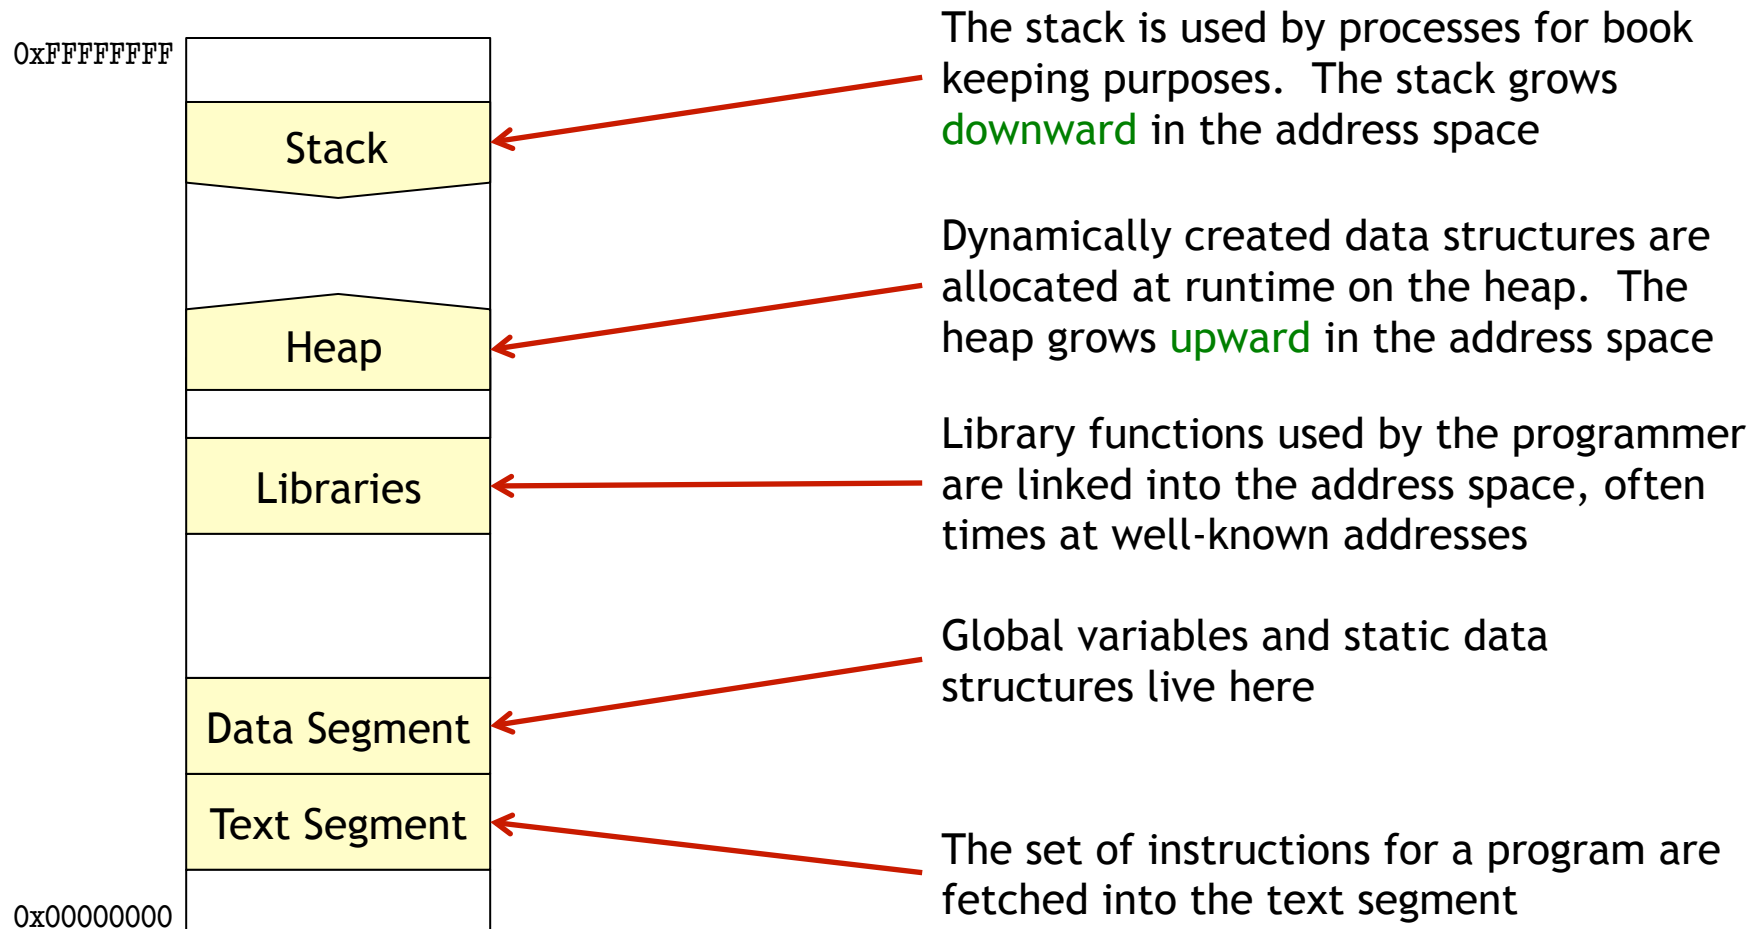
This attack vector is best used against root-level processes

- Why? The attacker's code will run as root!
- This provides a way for the attacker to “trick” the system into granting him privileged access



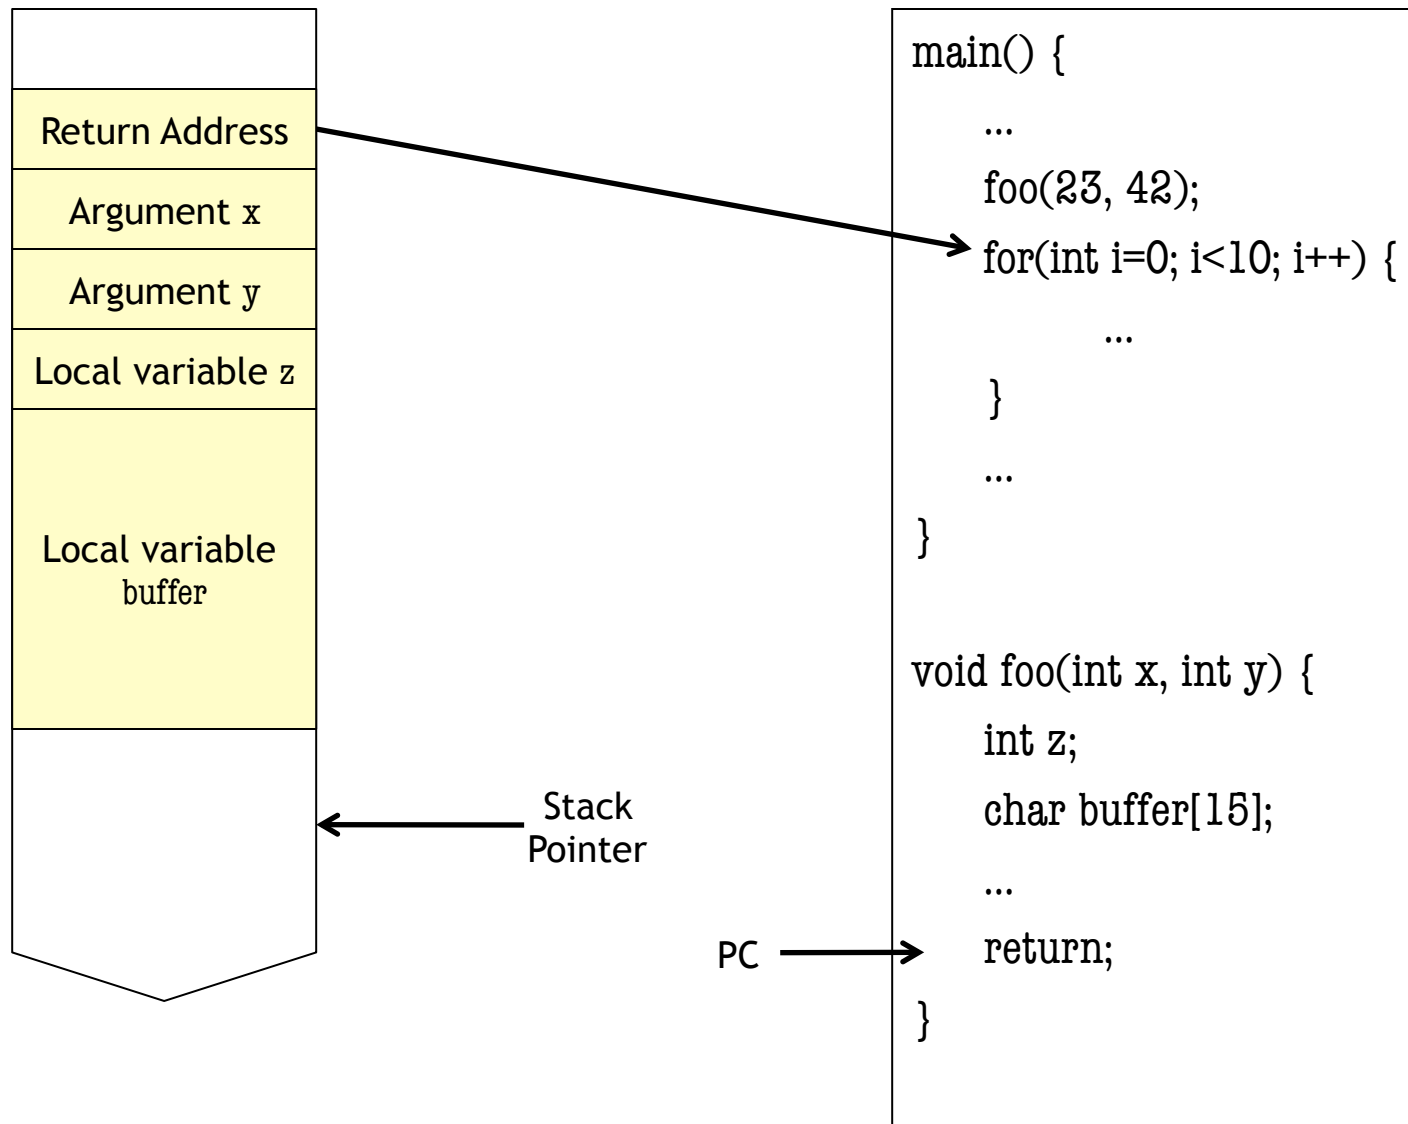
Understanding how buffer overflows work means understanding system memory management

Usually, a process's **address space** is a chunk of virtual memory accessed as an array of bytes





How does the stack work?





Smashing the stack

Aleph One, Smashing the Stack for Fun and Profit, Phrack Issue 49, Nov. 1996.

In a stack smashing attack, the attacker's goal is to:

- Inject exploit code onto the stack
- Overwrite the return address pointer to transfer control into the exploit code, rather than the calling function

Result: If the exploited code was running as an administrator the attacker's code executes with administrative privileges!

- Unix servers often run as root
- Most Windows users log in with administrative privileges




This can often be accomplished using unsafe string operators in C

For example, consider the following program fragment



```
void bar(char *s) {  
    char buffer[100];  
    strcpy(buffer,s);  
    ...  
}
```

Presumably, this returns
to main() or some other
calling function





How can we craft an exploit for the overflow?

Often our exploit code is just a few system calls

Caveat: Writing good shellcode is a bit of an art

- The attacker needs a good understanding of memory layout
- Code needs to be location independent
- If the code is exploiting C string functions, it cannot include any 0x00 bytes

Assuming that we can do that (see MetaSploit project) we must still answer two critical questions:

1. Where is the buffer and how big is it?
2. Where is the return address?

In reality, we only need to have **approximate** answers to these questions!



NOP sleds help us guess the program's start location on the stack

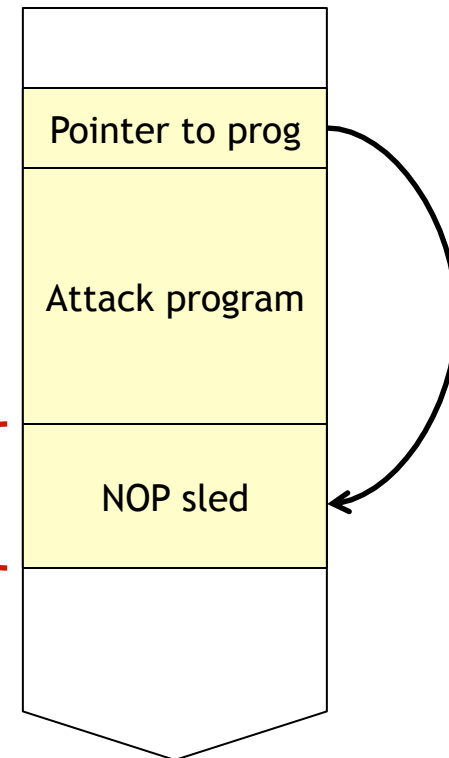
Most CPUs support a *No Operation* instruction of some sort

Rather than starting our attack program at the beginning of our buffer, we can pad it out with a bunch of NOP instructions

This is **good**, because the overwritten return address does not need to point exactly to the location of our exploit code!

We can point anywhere in here

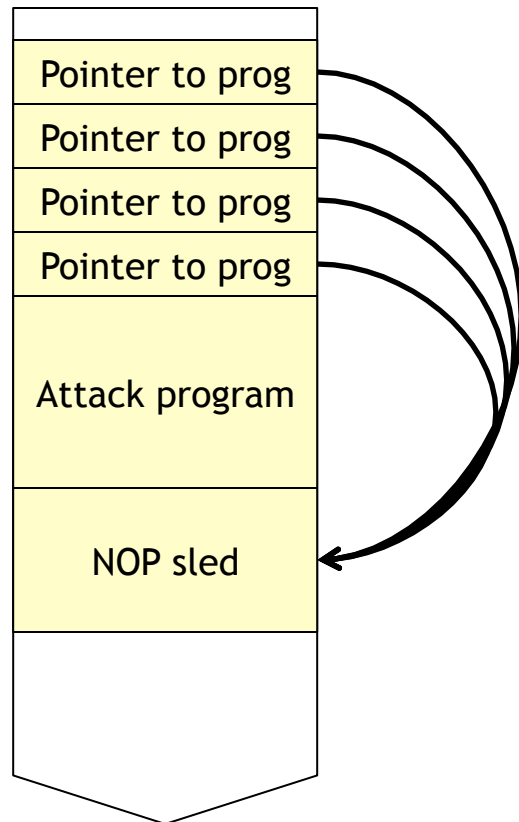
Unfortunately, we still need to know the **exact location of the return address**, which can be tricky...





We can also guess the location of the return address on the stack

How? Simply inject **several** new return addresses onto the stack!



As long as **one** of these pointers overwrites the actual return address location of the stack **and** it points back into the NOP sled, the attack will succeed



So... How do we defend against buffer overflows?

One way to stop buffer overflows is to use **non-executable stacks**

- If nothing can be executed from the stack, the exploit will fail
- Unfortunately, this is hard to implement in practice, as legacy compilers tend to like putting executable code in the stack
- A variation of this approach involves making writable memory pages non-executable

Another defense mechanism is the use of **canary words**

- A canary word is a random word of data written just before the return address on the stack
- Odds are, if the return address is overwritten, so is the canary word
- See Crispin Cowan's *StackGuard* for more information





So... How do we defend against buffer overflows?

Some researchers have proposed the use of **address space randomization** (ASR) to protect against stack smashing [1,2]

- **Intuition:** Shuffle the locations of variables and static data stored on the stack randomly each time the program is executed
- If the buffer is overflowed, there's really no way* to tell what data structures will be overwritten
- Unfortunately, researchers have shown that this certain variations of this approach can be defeated [3]

Another way to stop stack smashing is to use a type-safe programming language (e.g., Java or C#) ☺

[1] S. Bhatkar, R. Sekar, and D. DuVarney. Efficient Techniques for Comprehensive Protection from Memory Error Exploits," 14th USENIX Security Symposium, August 2005.

[2] <http://pax.grsecurity.net>

[3] H. Shacham, M. Page, B. Pfaff, E.-J. Goh, N. Modadugu, and D. Bohen, "On the Effectiveness of Address-Space Randomization," 11th ACM Conference on Computer and Communications Security, pages 298-307, Oct. 2004.



Stack smashing has a number of variants

If an attacker cannot launch a “traditional” stack smashing attack, there are other avenues that she could follow

Heap smashing is like stack smashing, but corrupts heap data

- There is much more variety across systems with how the heap is managed
- Attacks are difficult to construct, and typically are specific to a specific machine and/or configuration

Return to libc attacks jump to existing library code, not attacker code

- The attacker first correctly places her arguments on the stack
- Then she sets the return pointer to jump to pre-loaded libraries (e.g., libc)

Return oriented programming generalizes return to libc

- Generate a Turing complete library of computational “widgets” from the libc library (or any other library)
- Inject an attack vector that puts widgets together in some desirable way
- The effect is running arbitrary code **without** injecting executable code!

E. Buchanan, R. Roemer, H. Shacham, and S. Savage, “When Good Instructions Go Bad: Generalizing Return-Oriented Programming to RISC,” CCS 2008.



Buffer Overflow Wrap-Up

In a stack smashing attack, the attacker alters kernel data structures to make a process behave in an unexpected manner

If the process that is attacked is running with elevated privileges (e.g., a server running as root), then the attacker can do all sorts of fun things

- Set up tunnels into the system
- Alter the password file
- Install rootkits to gain long term control of the system
- Delete files
- ...

Clearly, this constitutes a violation of integrity of mechanism

What about violating constraints within an application?



Sometimes, we don't need to subvert the entire OS to wreak havoc.

For example, many times applications have their own security mechanisms in place

- These constraints control the functionality exposed to different users
- Enforced by the application, not the OS
- In this case, compromising the application is all that is needed

Case Study: SQL Injection



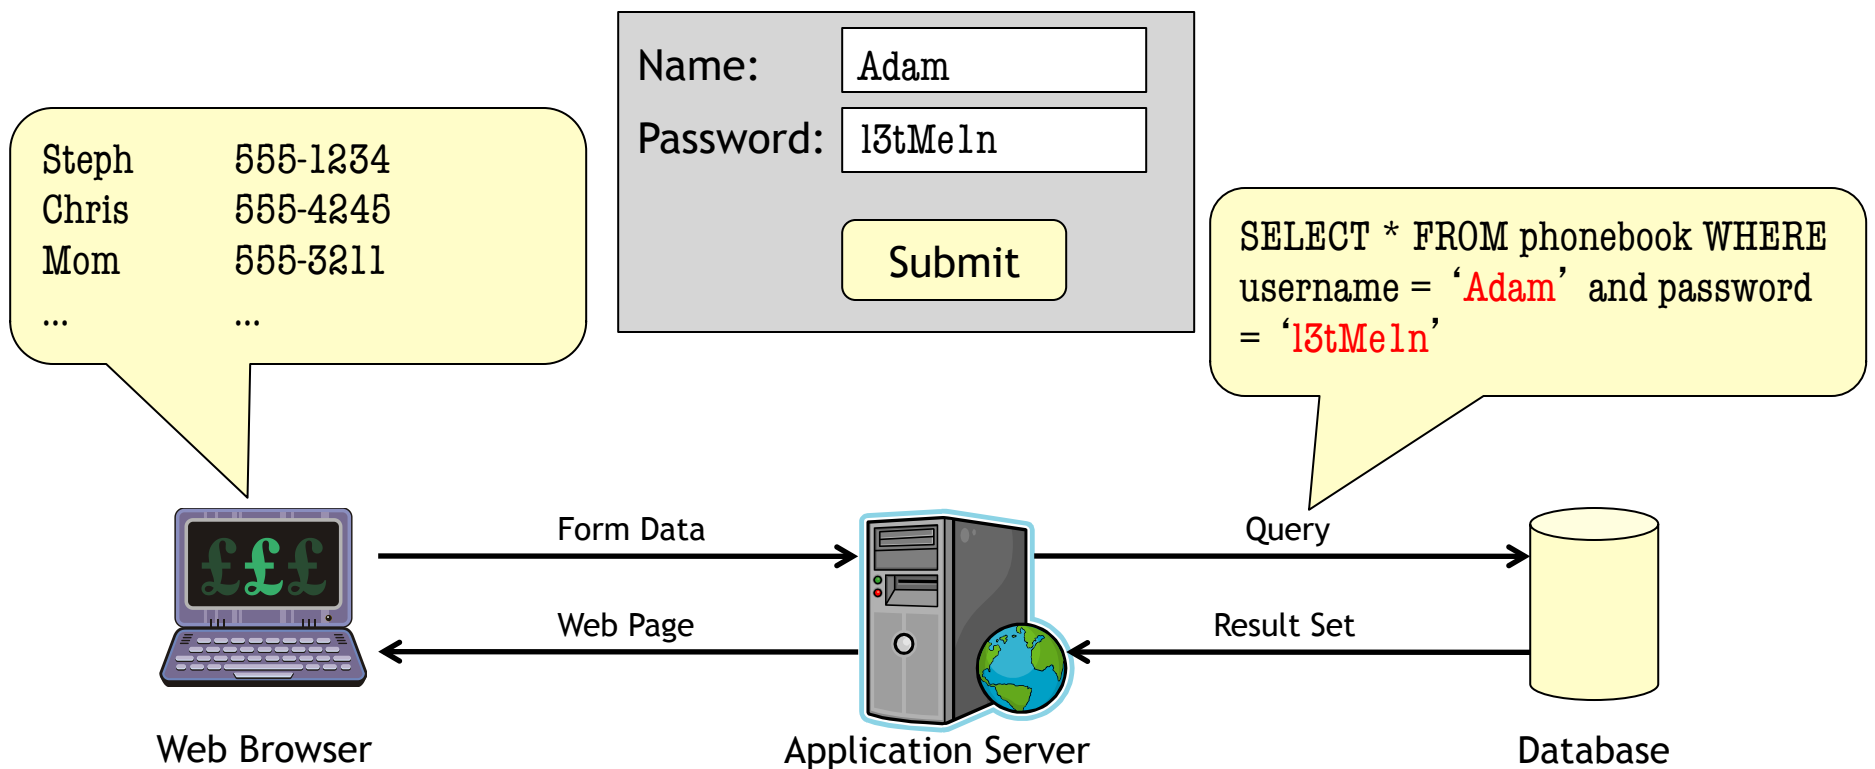
<http://xkcd.com/327/>

SQL injection attacks are specific to application servers with database back ends

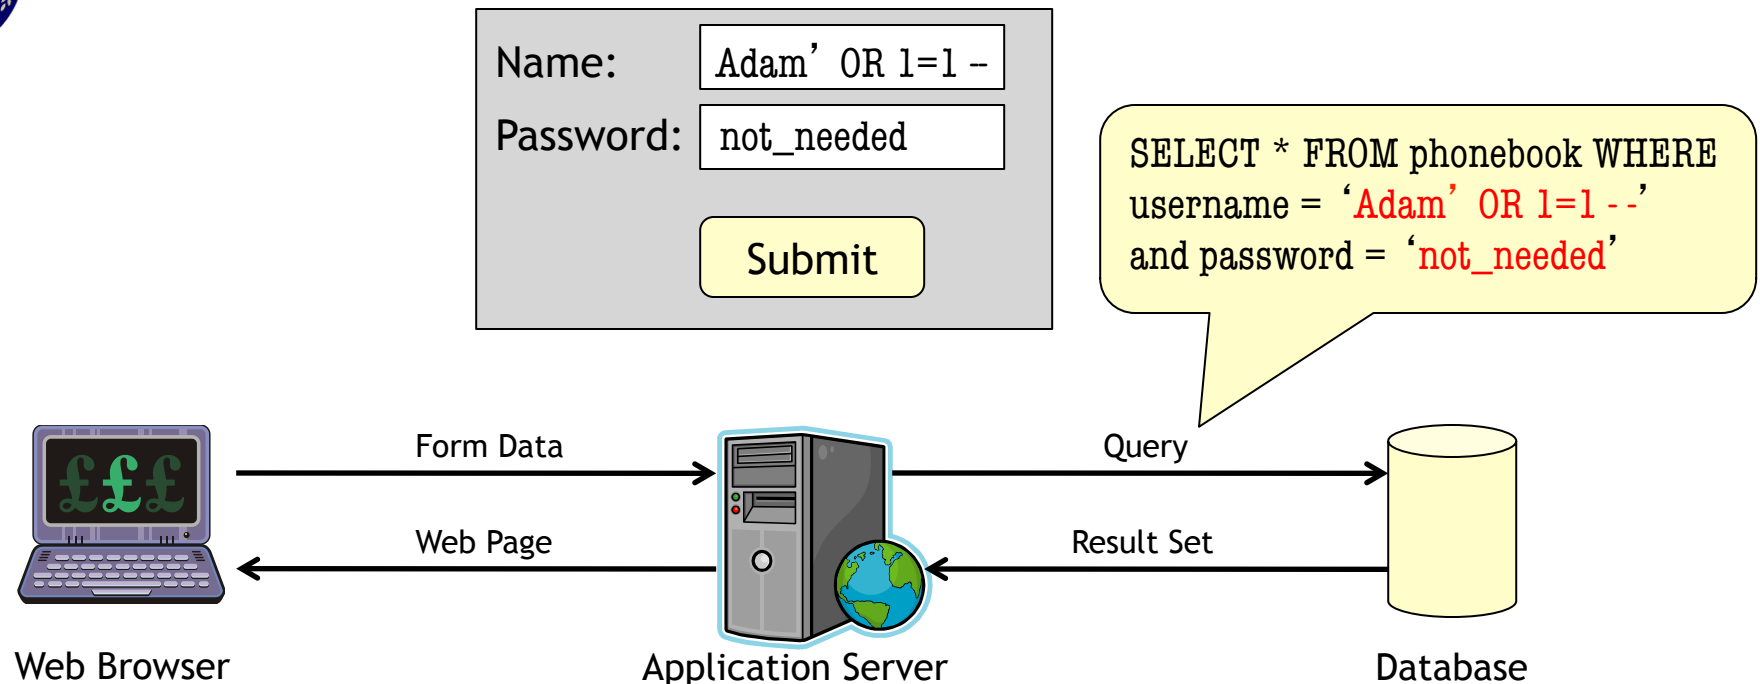


These applications serve **dynamic content** based upon queries that are parameterized by user input

Example: An online phone book application



What could possibly go wrong with such a simple application?



Note that:

- The test `1=1` is always true!
- Also, `--` denotes a comment in SQL, so the password check is bypassed!

Result: All rows are displayed!




SQL injection can be used for all kinds of attacks...




Gaining access to the system

- Perform an INSERT on the “Users” table



Modify data in the database

- Drop tables
- Change salaries
- Etc.



Perform privilege escalation attacks

- Stealing passwords
- Inserting rights
- Etc.

Clearly, this is much easier than crafting a buffer overflow!



The real difficulty lies in ascertaining the structure of the target database, but even this is not so difficult... So how we defend ourselves?

One solution is to sanitize database inputs

- For example, ' → \', " → \", and \ → \\
- Unfortunately, this doesn't always work...

➤ SELECT * FROM salary WHERE id = 23 OR 1=1

Check input field validity

- Fields typically have fixed vocabularies
- Make sure illegal characters don't appear
- Unfortunately, this isn't always possible

Place length limits on the allowed inputs to form fields

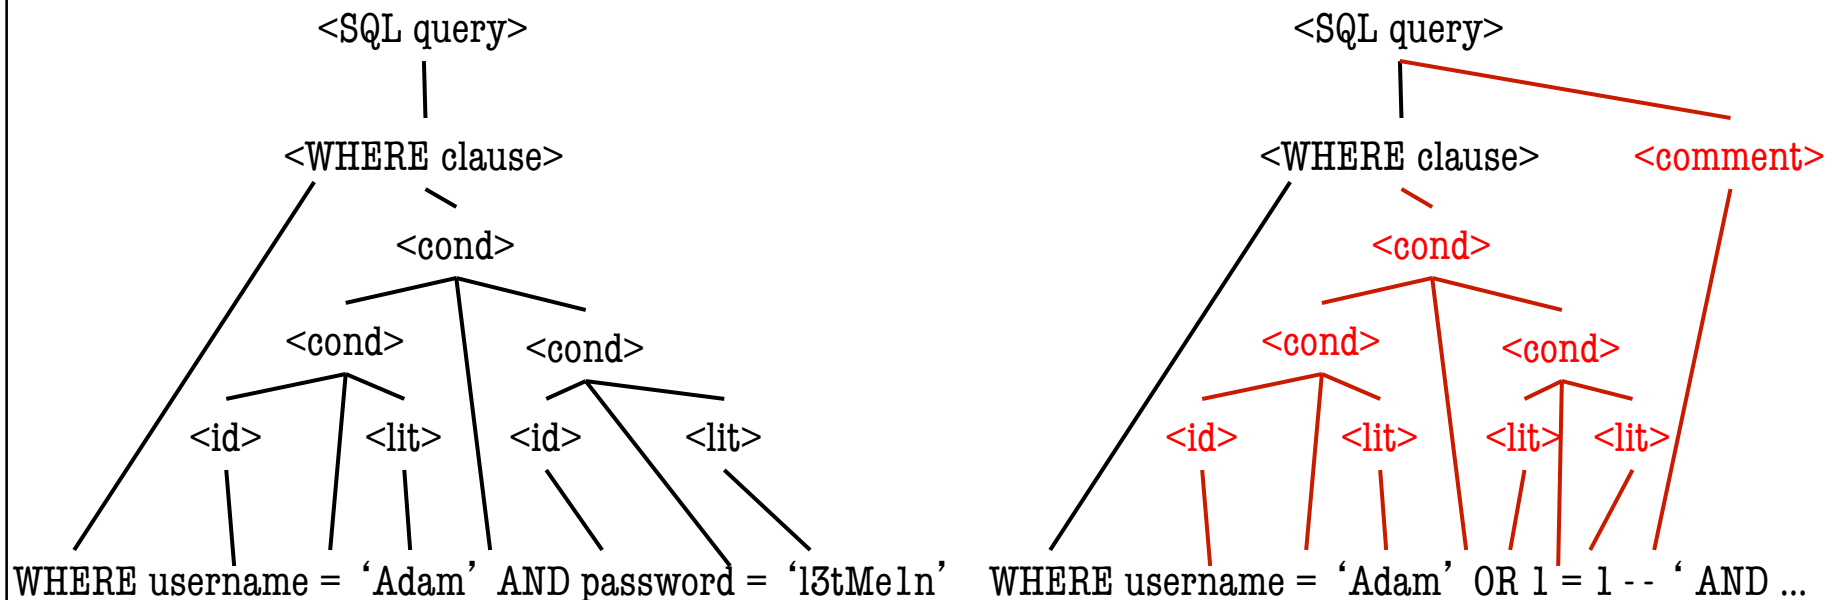
These are all basically hacks...



A solution proposed by Bandhakavi et al. uses parse tree comparison to detect SQL injection

Intuition: The parse tree for an SQL injection query is almost always different than the parse tree for the programmer intended query

Example:



Sruthi Bandhakavi, Prithvi Bisht, Madhusudan Parthasarathy, V.N. Venkatakrishnan, "CANDID: Preventing SQL Injection Attacks using Dynamic Candidate Evaluations," 14th ACM Conference on Computer and Communications Security (CCS), November 2007.



The problem that must be addressed is how to infer the programmer intended parse tree for a given query

Bandhakavi et al. solve this problem by generating benign **candidate inputs** for each dynamically generated query

The problem of **how** to generate inputs that are benign and do not alter the program's control flow is undecidable in general...

However, simple heuristics make this possible!

- Replace all integer inputs with "1"
- Replace all string inputs with an equal length string containing only "a"s
- Replace all booleans with "true"
- Make control flow decisions based upon **actual** input only


Parse trees are generated for the actual query and the candidate query

- Actual query is issued only if both parse trees match

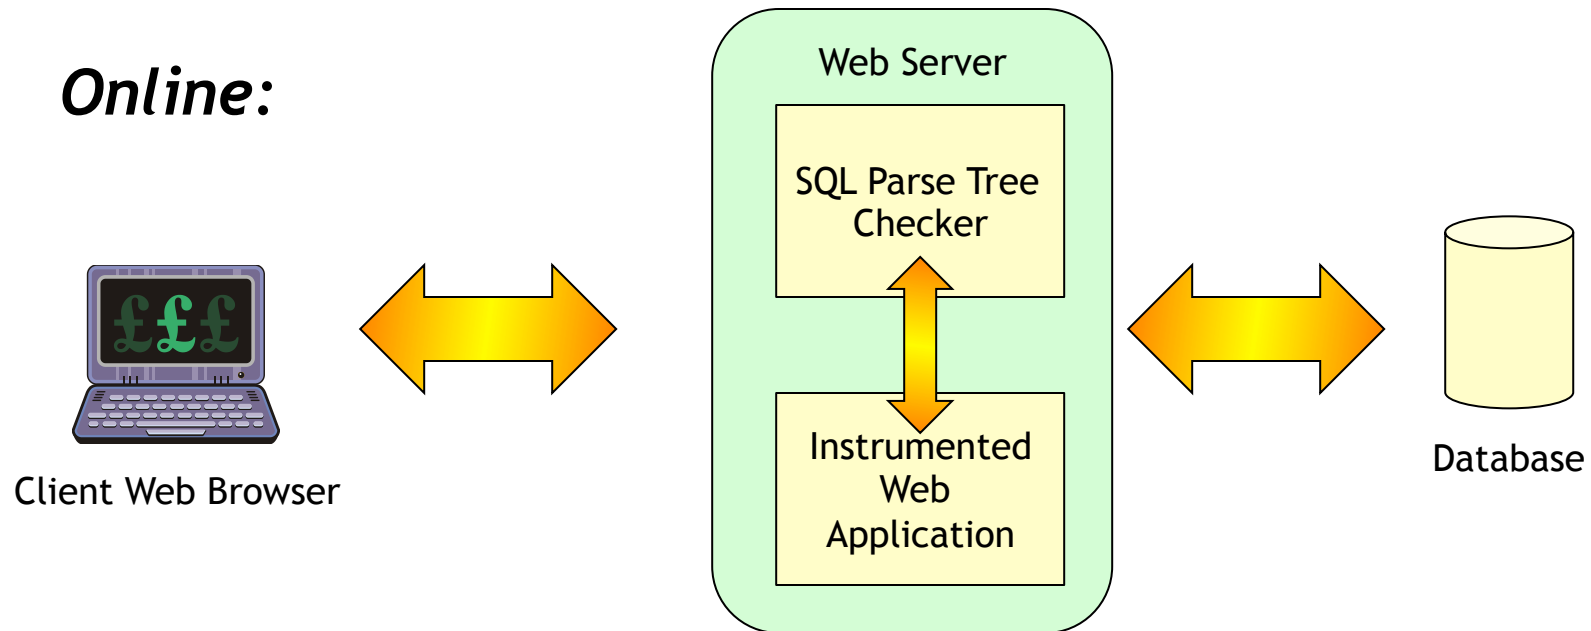


How is the CANDID system set up?

Offline:



Online:




Bandhakavi's solution performs well on standard database benchmarks



| Application | Lines Of Code | Successful Attacks / Attacks Detected | Non-attacks / False Positives | Parse Errors Expected/ Detected |
|-------------|---------------|---------------------------------------|-------------------------------|---------------------------------|
| Bookstore | 16959 | 1438 / 1438 | 2930 / 0 | 2124/2124 |
| EmployeeDir | 5658 | 1529 / 1529 | 2442 / 0 | 3067/3067 |
| Events | 7242 | 1414 / 1414 | 3320 / 0 | 2375/2375 |
| Classifieds | 10949 | 1475 / 1475 | 2076 / 0 | 3128/3128 |
| Portal | 16453 | 2995 / 2995 | 3415 / 0 | 1073/1073 |
| Checkers | 5421 | 262 / 262 | 7435 / 0 | 557/557 |
| Officetalk | 4543 | 390 / 390 | 2149 / 0 | 4060/4060 |

Bandhakavi's solution performs well on standard database benchmarks





Conclusions

OS Security relies on separating the **trusted** from the **untrusted**

The assumption of **integrity of mechanism** is central to security

Integrity of mechanism can be violated at the OS or application level

- **OS example:** Buffer overflow
- **Application example:** SQL injection

